

Idea de los Auto Encoders

Plan de Ataque

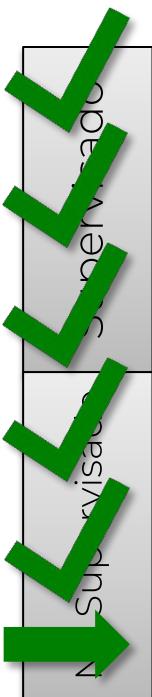
Plan de Ataque

Qué vamos a aprender en esta sección:

- Auto Codificadores (Auto Encoders)
- Cómo entrenar un Auto Encoder
- Sobocompletar las capas ocultas
 - Sparse Auto Encoders
 - Denoising Autoencoders
 - Contractive Autoencoders
- Stacked Autoencoders
- Deep Autoencoders

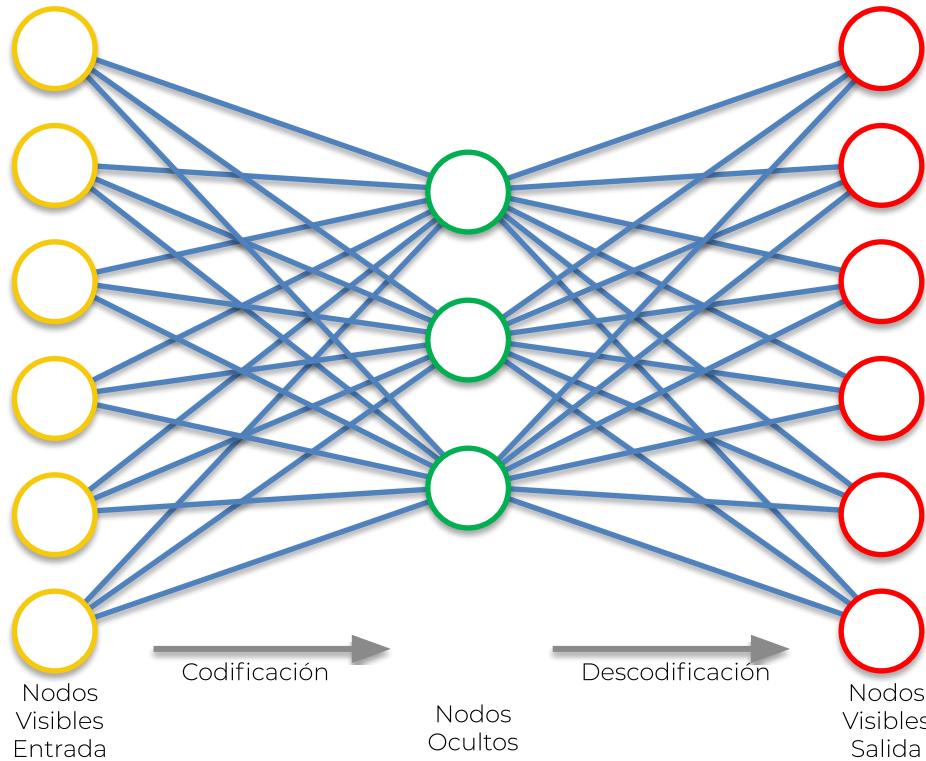
Auto Encoders

Auto Encoders

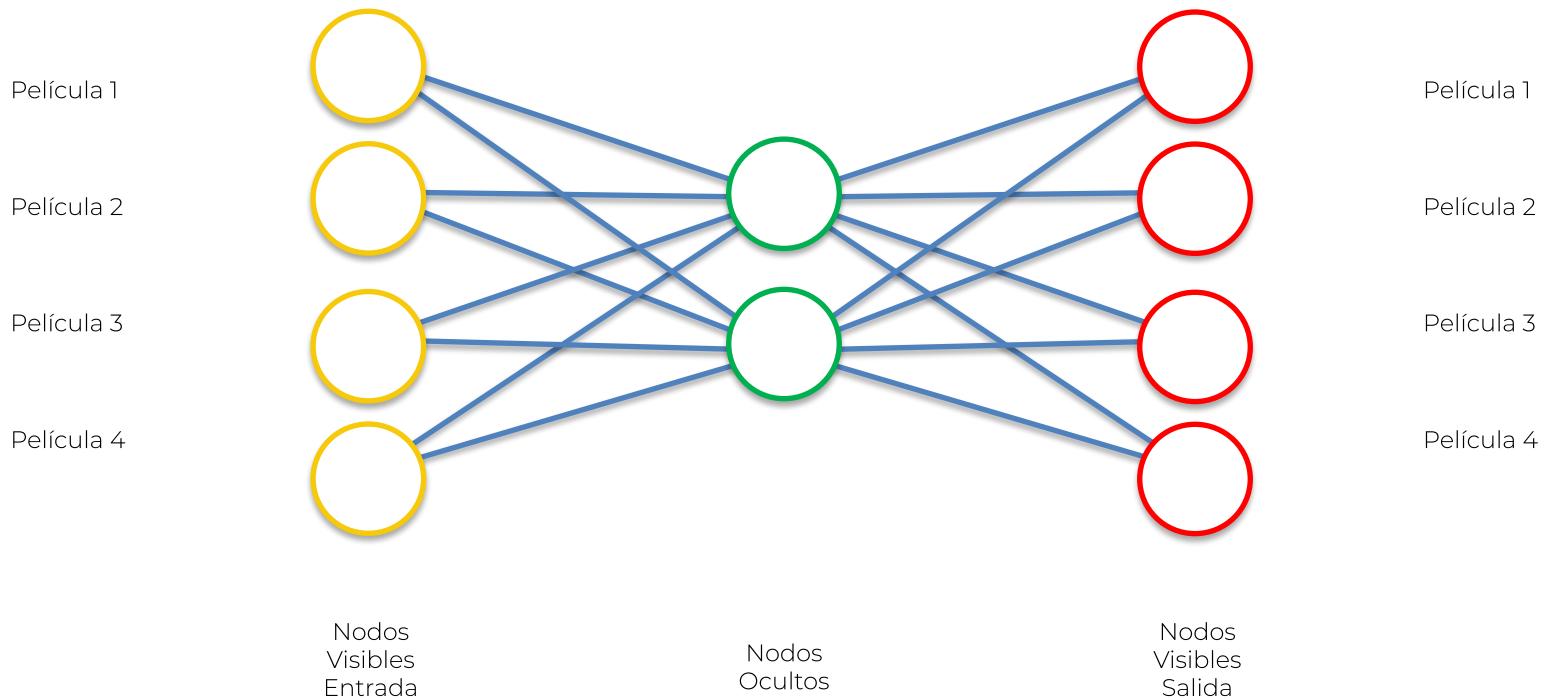


Redes neuronales artificiales	Utilizado para regresión y clasificación
Redes neuronales convolucionales	Utilizado para la visión por ordenador
Redes neuronales recurrentes	Utilizado para el análisis de series temporales
Mapas autoorganizados	Utilizado para la detección de características
Máquinas de Boltzmann Profundas	Utilizado para los sistemas de recomendación
AutoEncoders	Utilizado para los sistemas de recomendación

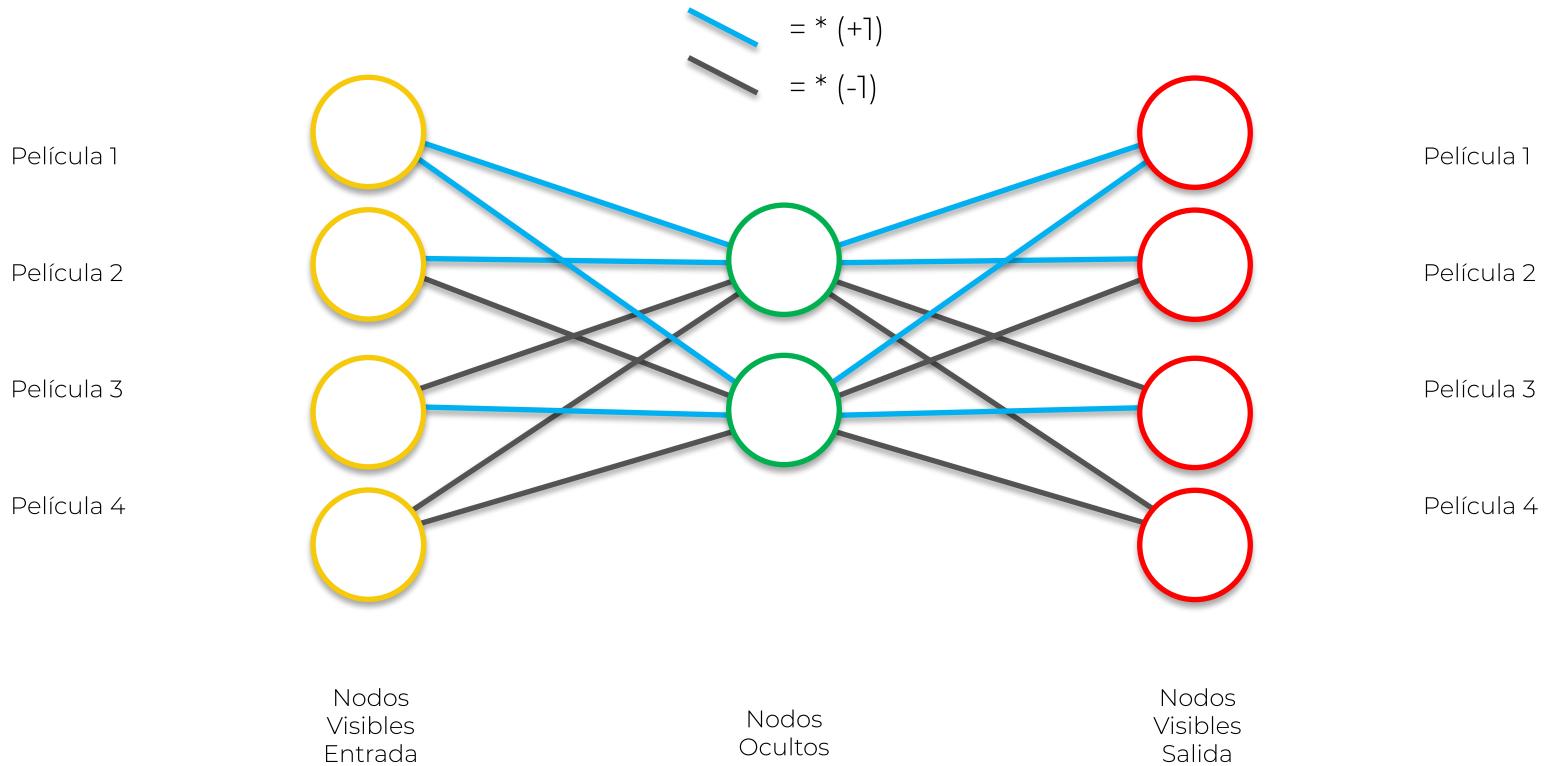
Auto Encoders



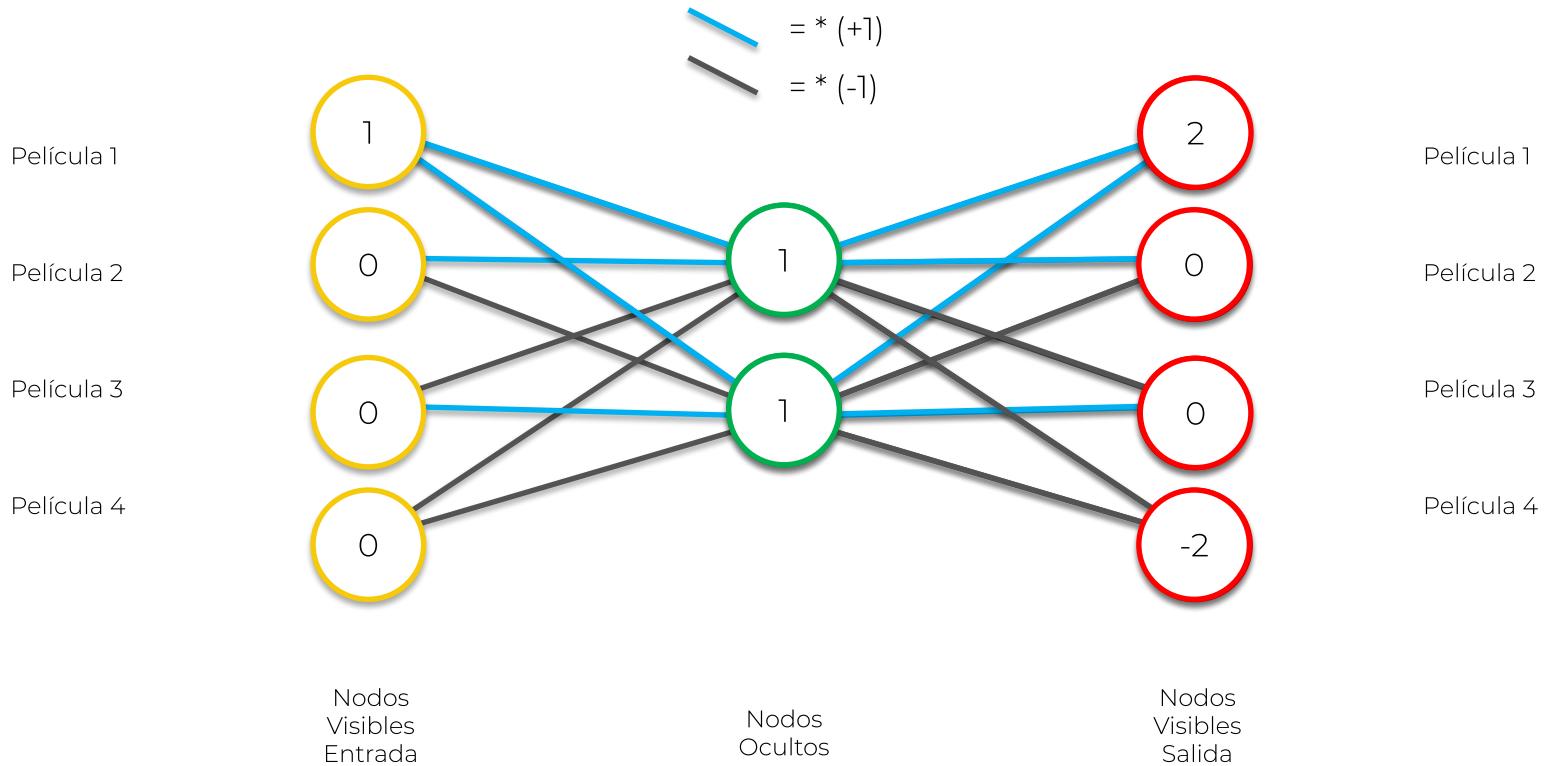
Auto Encoders



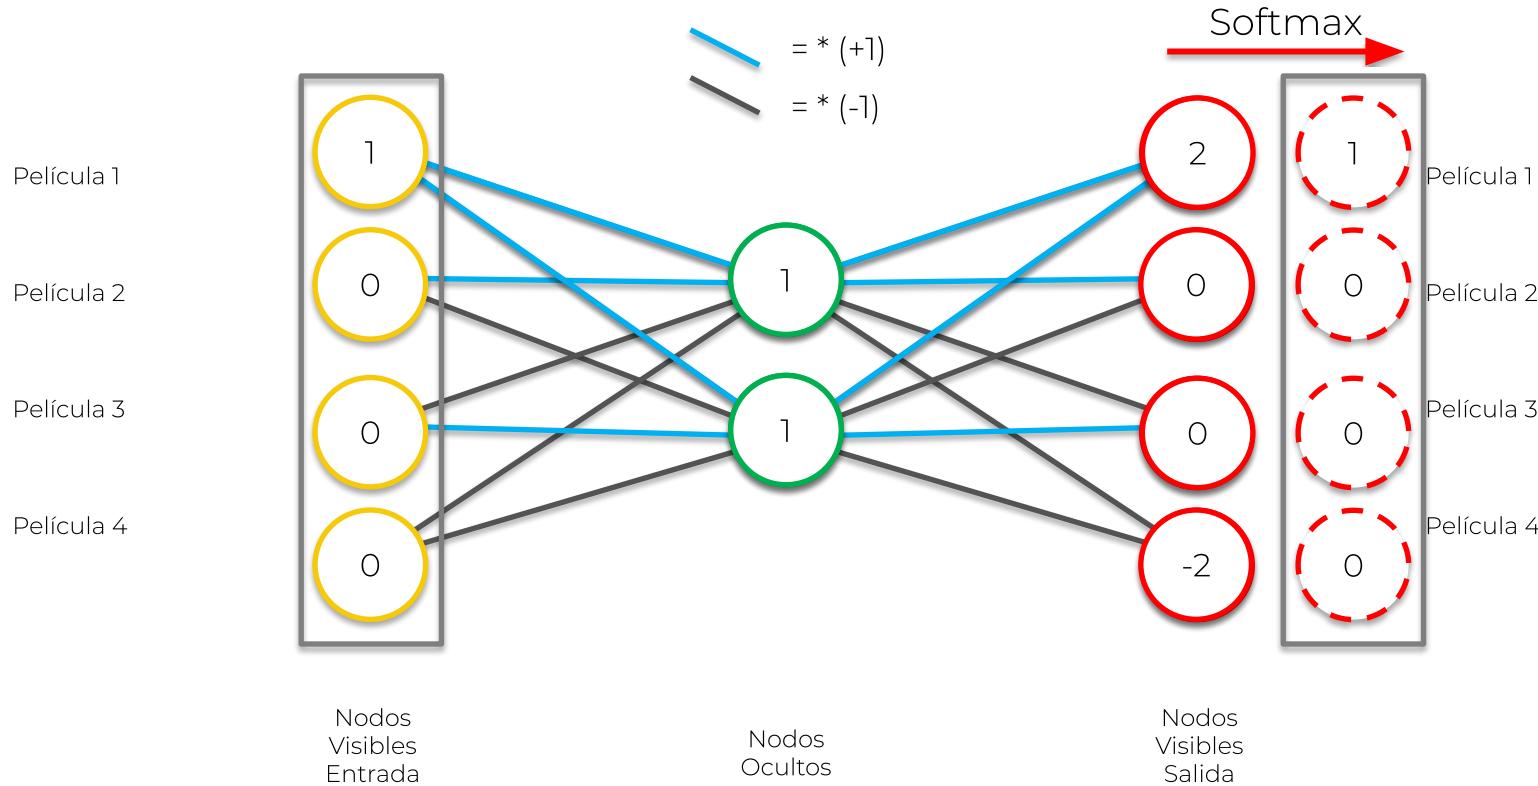
Auto Encoders



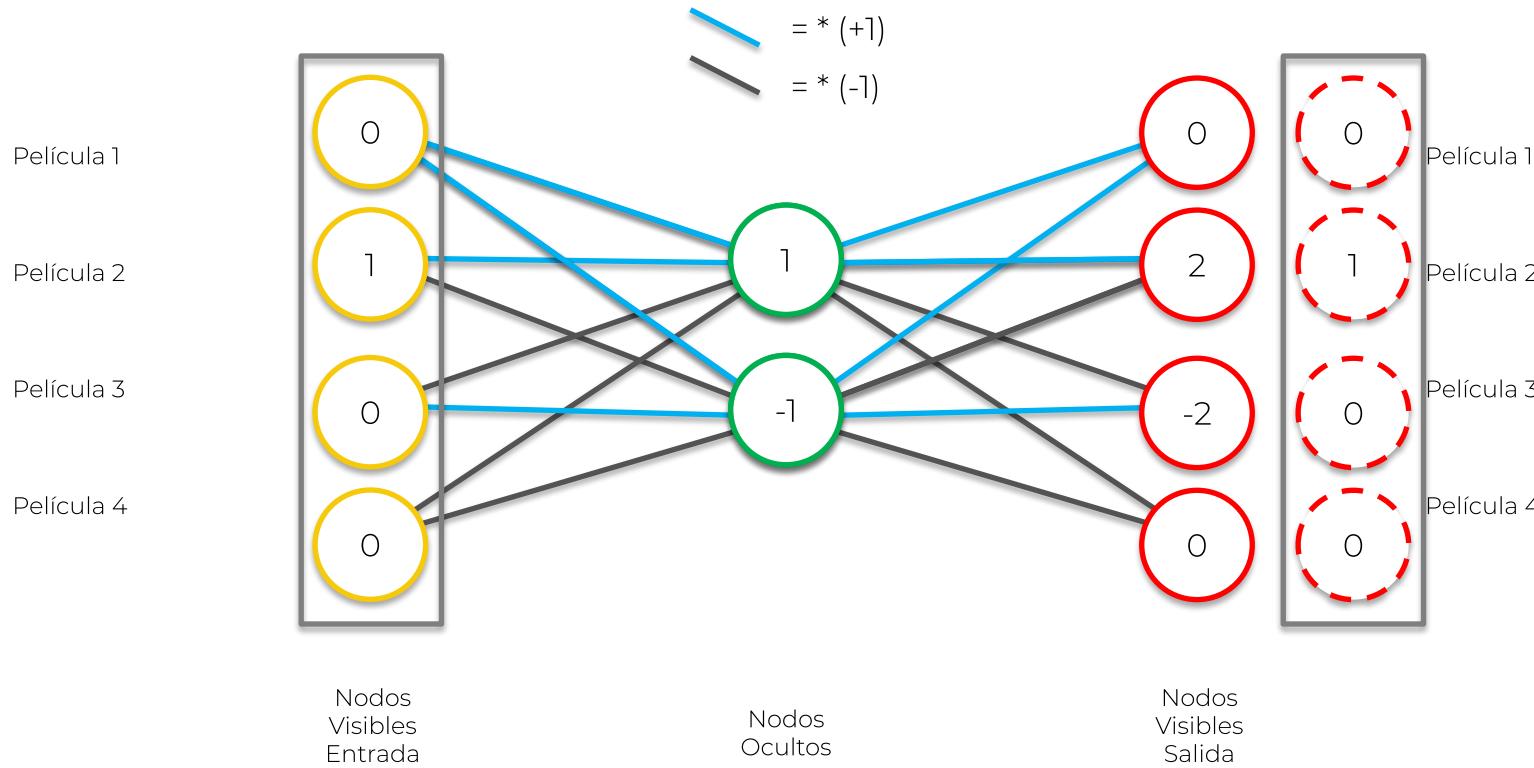
Auto Encoders



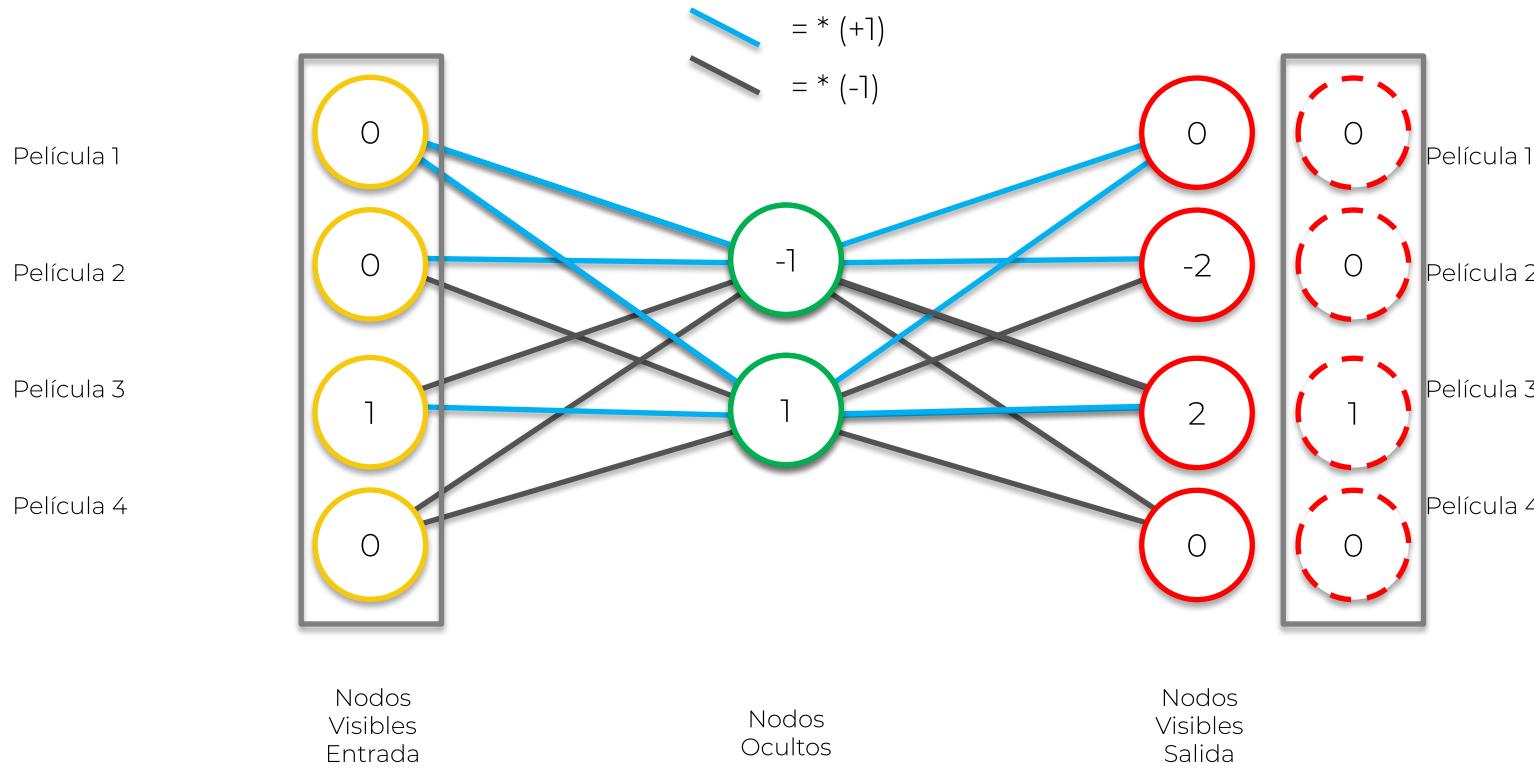
Auto Encoders



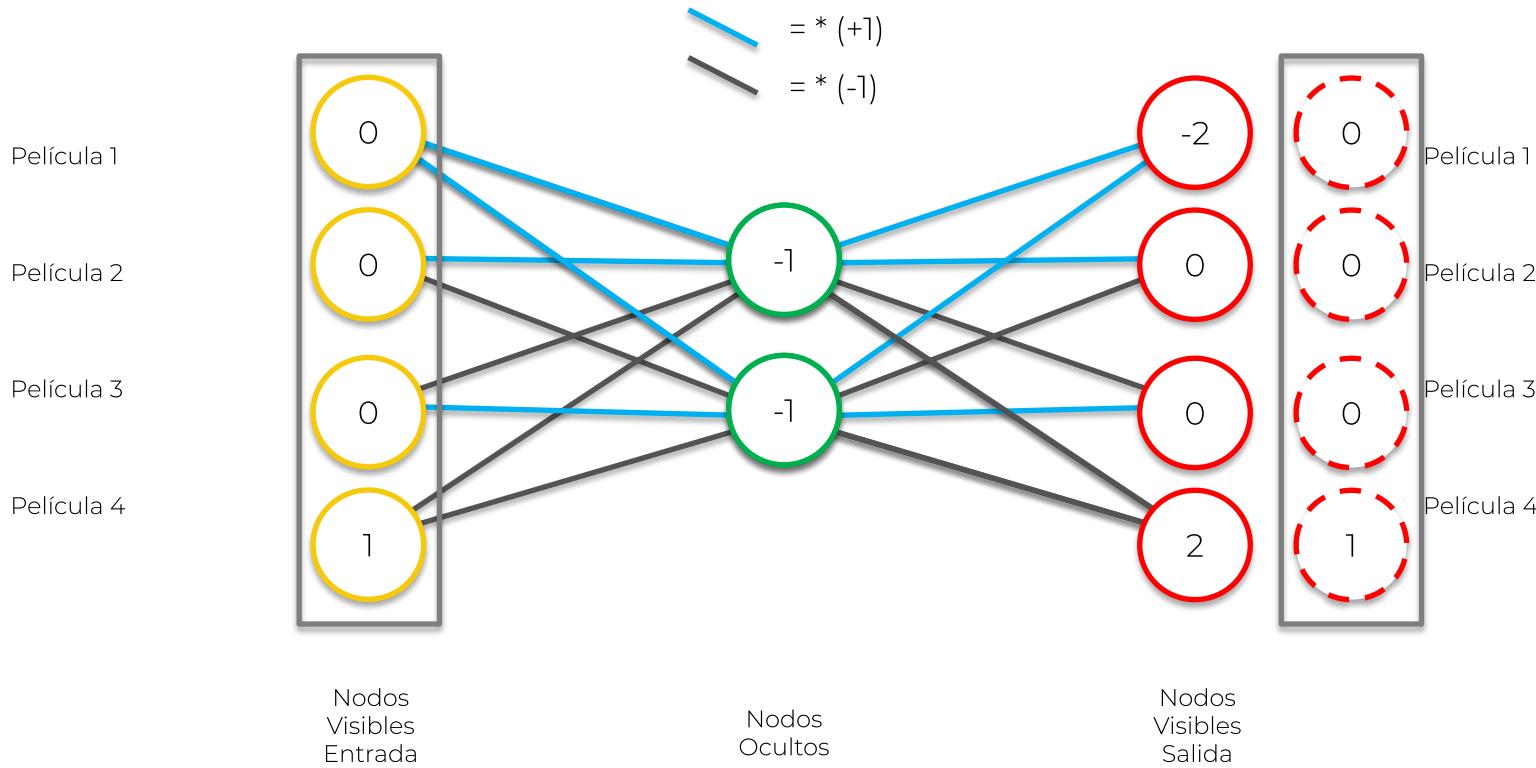
Auto Encoders



Auto Encoders



Auto Encoders



Auto Encoders

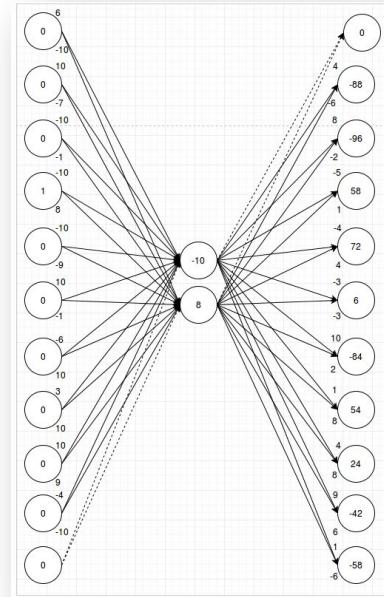
Lectura Adicional:

Neural Networks Are Impressively Good At Compression

de Malte Skarupke (2016)

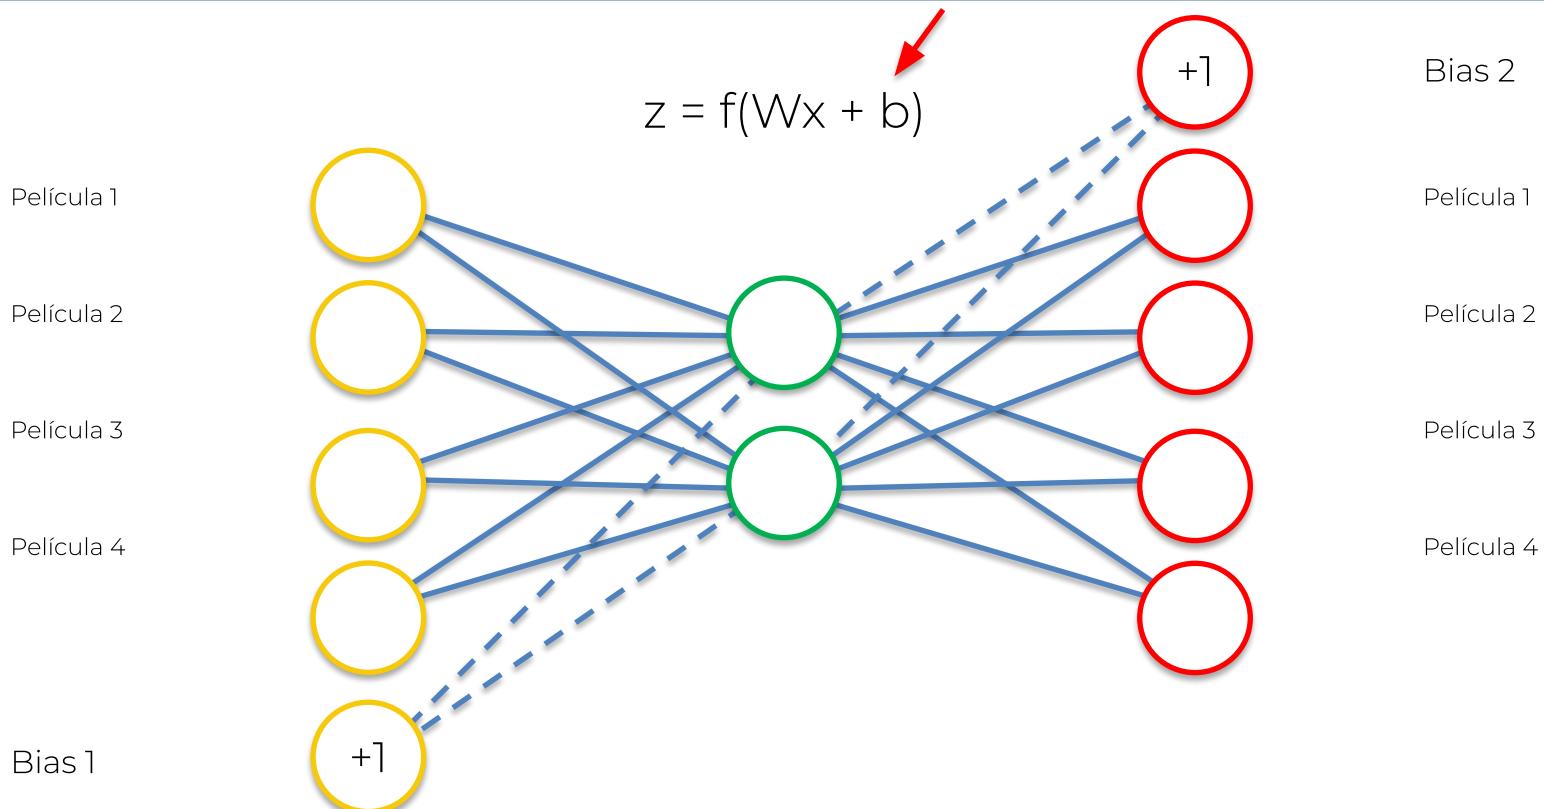
Link:

<https://probablydance.com/2016/04/30/neural-networks-are-impressively-good-at-compression/>

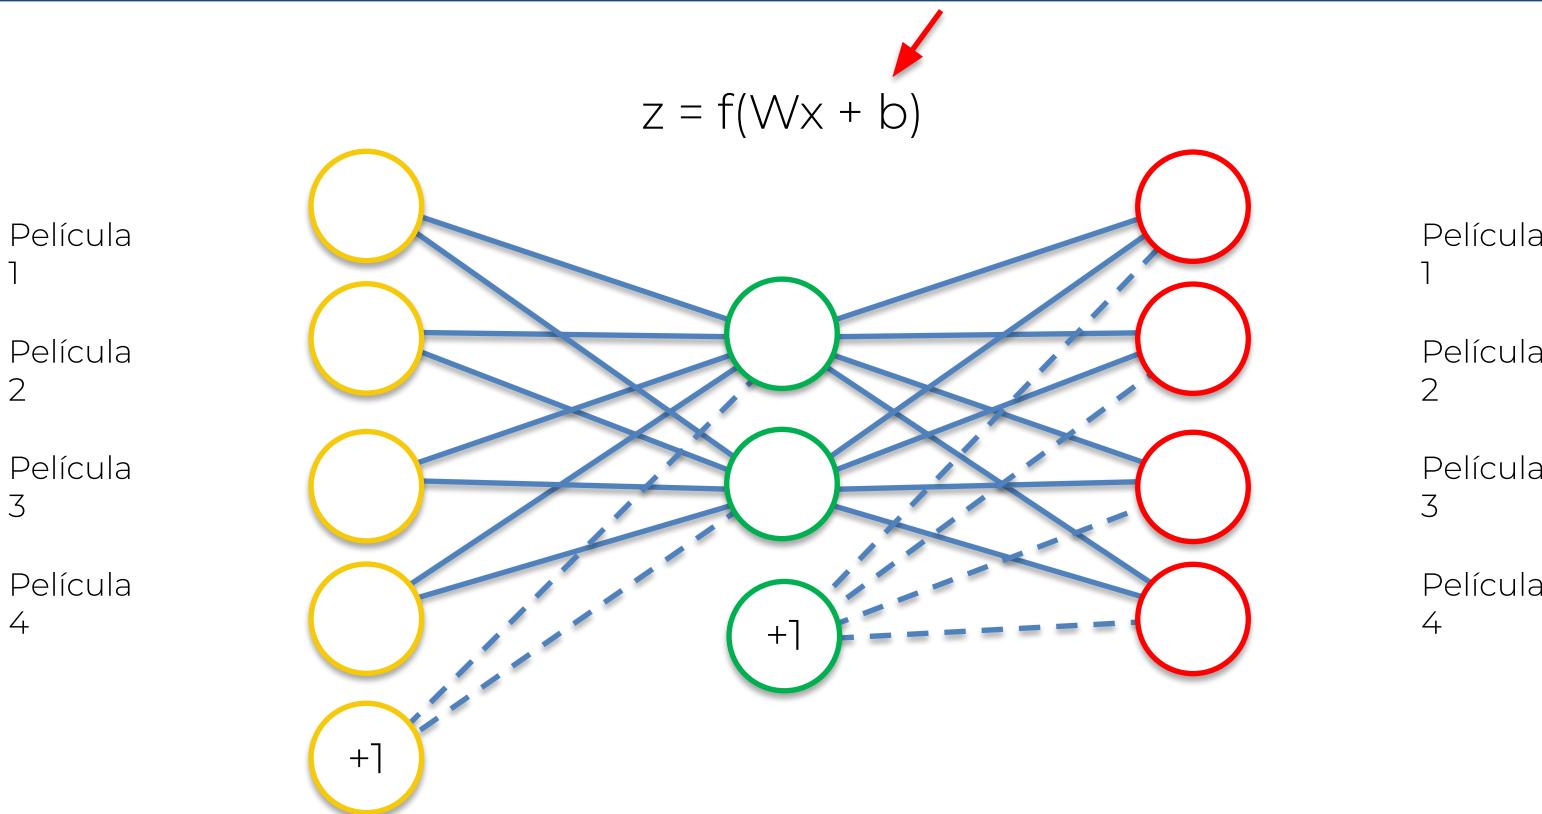


Acerca del término independiente (Bias)

Auto Encoders

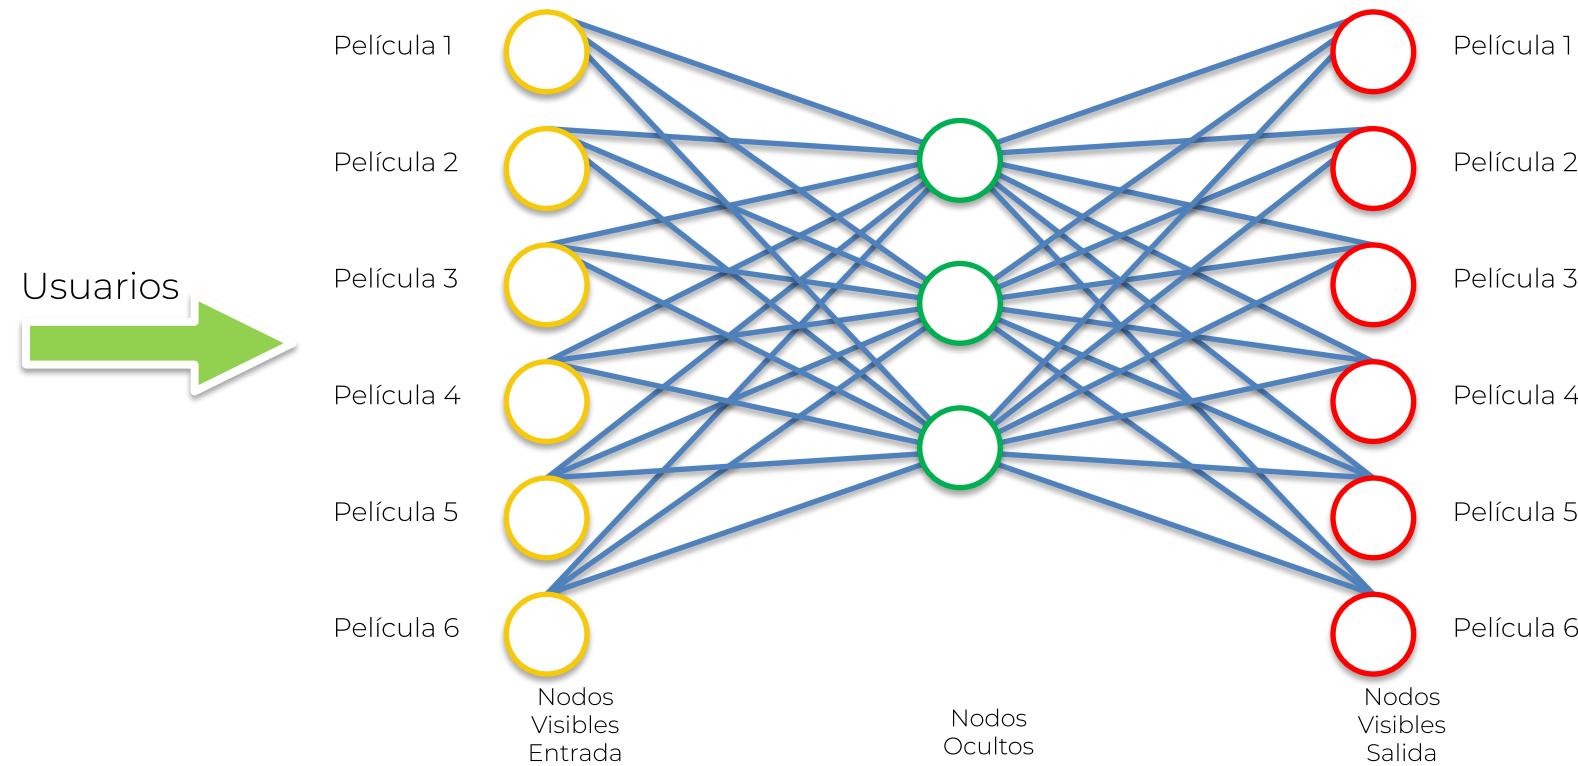


Auto Encoders



Cómo entrenar un Auto Encoder

Cómo entrenar un Auto Encoder



Cómo entrenar un Auto Encoder

PASO 1: Empezamos con un array donde las filas (observaciones) se corresponden a usuarios y las columnas (características) se corresponden a las películas. Cada celda (u, i) contiene la valoración (de 1 a 5, 0 si no hay valoración) de la película i valorada por el usuario u .



PASO 2: El primer usuario entra en la red. El vector de entrada $x = (r_1, r_2, \dots, r_m)$ tiene las valoraciones de todas las pelis.



PASO 3: El vector de entrada x se codifica en un vector z de menor dimensión con la función de transformación f (p.e: una función sigmoide):



$$z = f(Wx + b) \text{ donde } W \text{ es el vector de entrada de pesos y } b \text{ el bias}$$

PASO 4: z se descodifica en el vector de salida y de la misma dimensión que x para replicar el vector de entrada x .



PASO 5: Se calcula la reconstrucción del error $d(x, y) = ||x-y||$ con el objetivo de minimizarlo.



PASO 6: Back-Propagation: de derecha a izquierda se propaga el error. Los pesos se actualizan según su responsabilidad en el error. El ratio de error decide cuánto se actualizan dichos pesos.



PASO 7: Se repiten los pasos 1 a 6 y se actualizan los pesos tras cada observación (Reinforcement Learning). O:
Se repite los pasos 1 a 6 pero se actualizan los pesos solo tras un conjunto de observaciones(Batch Learning).



PASO 8: Cuando todo el conjunto de entrenamiento ha pasado por la RNA se completa un epoch. Repetir más veces.

Cómo entrenar un Auto Encoder

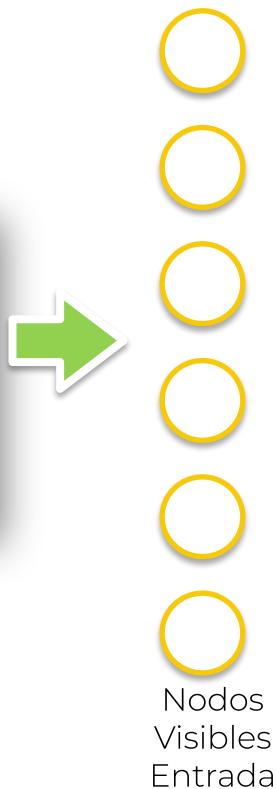
PASO 1

	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6
User 1	1	0		1	1	1
User 2	0	1	0	0	1	0
User 3		1	1	0	0	
User 4	1	0	1	1	0	1
User 5	0		1	1		1
User 6	0	0	0	0	1	
User 7	1	0	1	1	0	1
User 8	0	1	1		0	1
User 9		0	1	1	1	1
User 10	1		0	0		0
User 11	0	1	1	1	0	1

Cómo entrenar un Auto Encoder

PASO 2

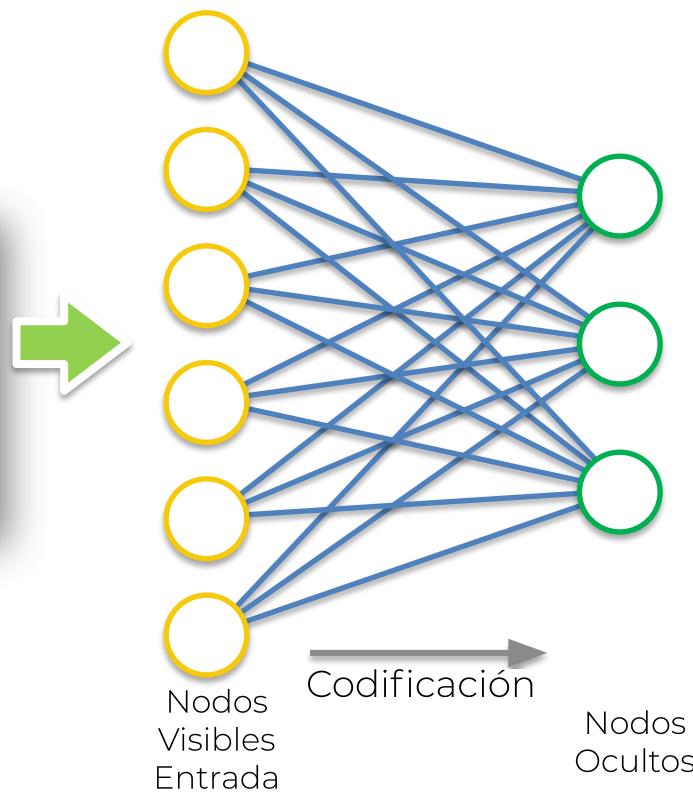
	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6
User 1	1	0		1	1	1
User 2	0	1	0	0	1	0
User 3		1	1	0	0	
User 4	1	0	1	1	0	1
User 5	0		1	1		1
User 6	0	0	0	0	1	
User 7	1	0	1	1	0	1
User 8	0	1	1		0	1
User 9		0	1	1	1	1
User 10	1		0	0		0
User 11	0	1	1	1	0	1



Cómo entrenar un Auto Encoder

PASO 3

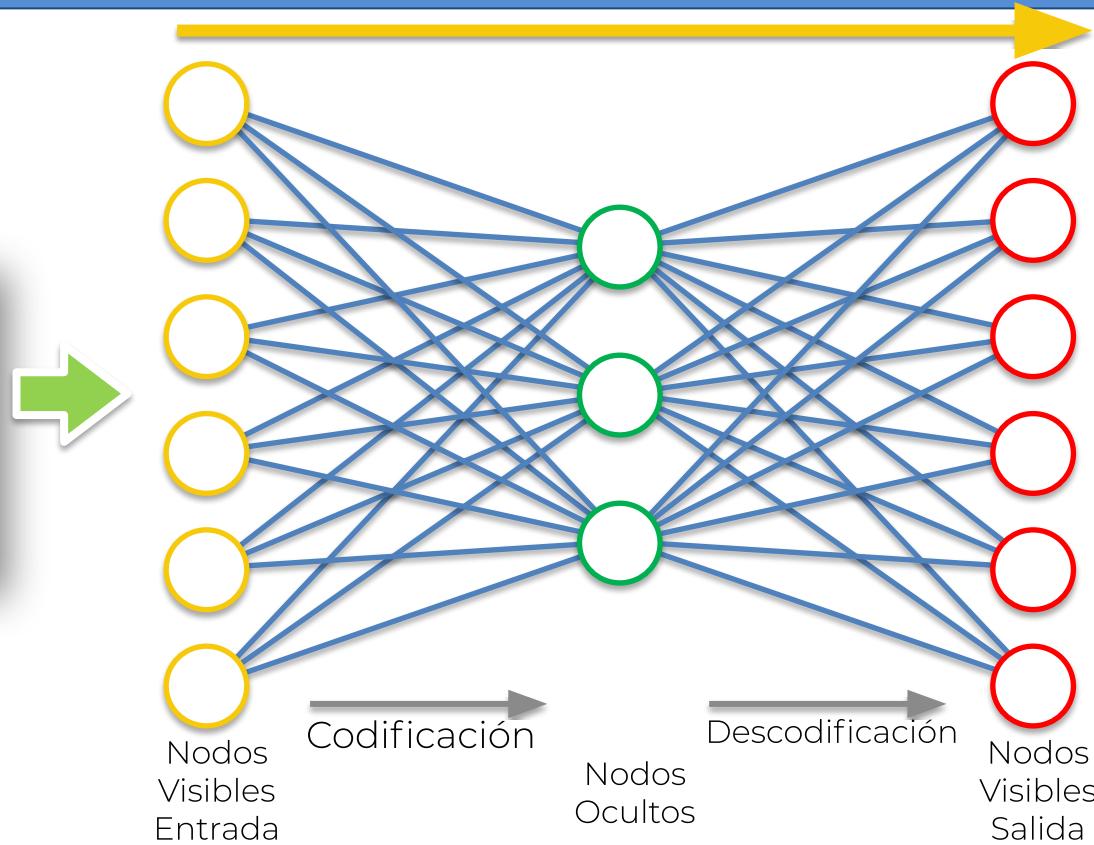
	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6
User 1	1	0				
User 2	0	1	0	0	1	0
User 3		1	1	0	0	
User 4	1	0	1	1	0	1
User 5	0		1	1		1
User 6	0	0	0	0	1	
User 7	1	0	1	1	0	1
User 8	0	1	1		0	1
User 9		0	1	1	1	1
User 10	1		0	0		0
User 11	0	1	1	1	0	1



Cómo entrenar un Auto Encoder

PASO 4

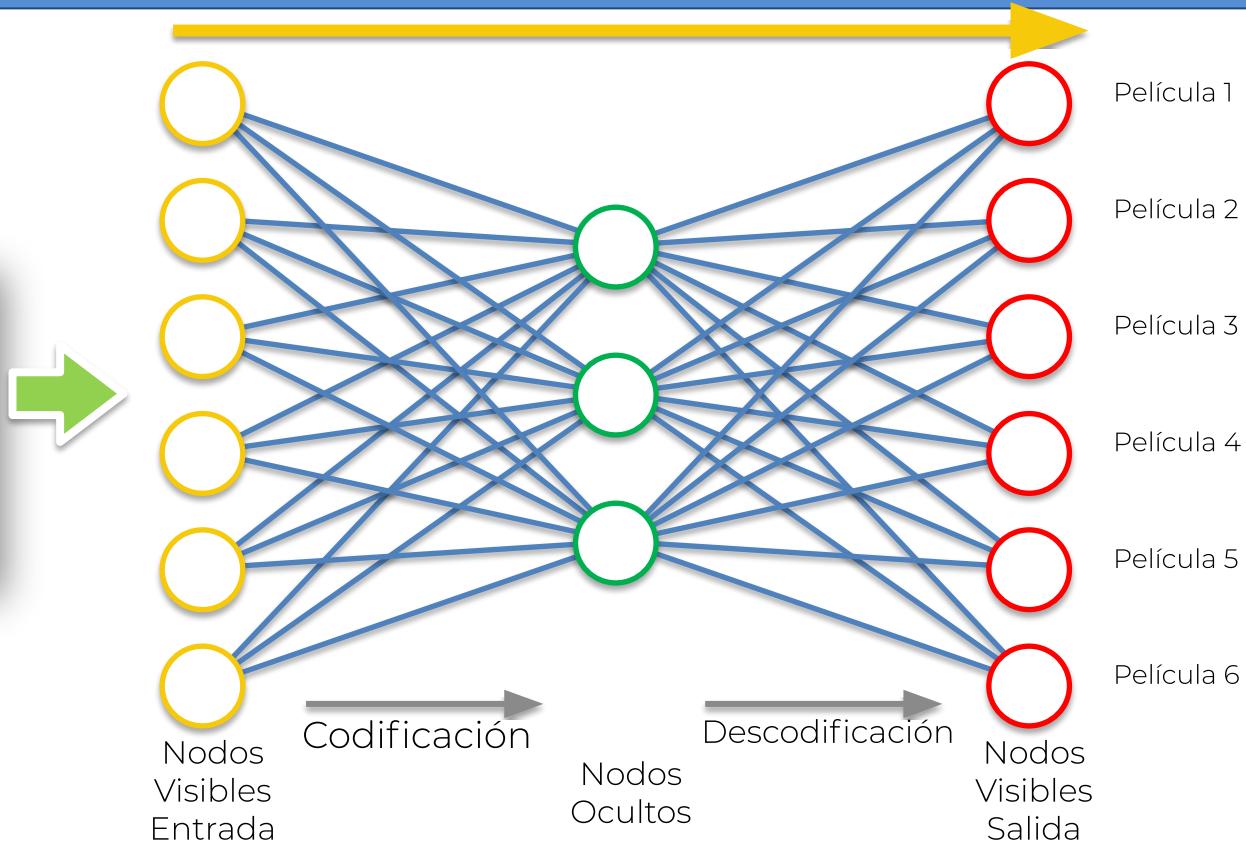
	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6
User 1	1	0				
User 2	0	1	0	0	1	0
User 3		1	1	0	0	
User 4	1	0	1	1	0	1
User 5	0		1	1		1
User 6	0	0	0	0	1	
User 7	1	0	1	1	0	1
User 8	0	1	1		0	1
User 9		0	1	1	1	1
User 10	1		0	0		0
User 11	0	1	1	1	0	1



Cómo entrenar un Auto Encoder

PASO 5

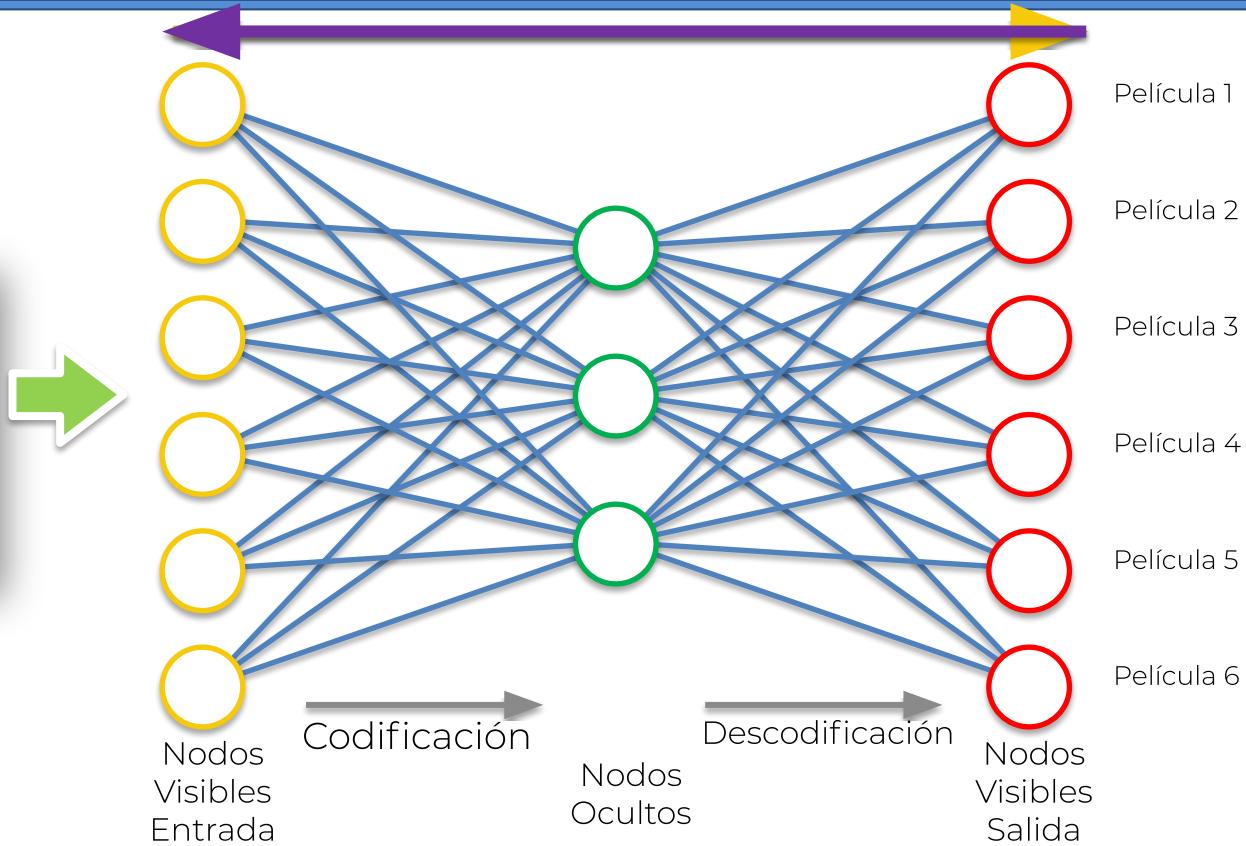
	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6
User 1	1	0				
User 2	0	1	0	0	1	0
User 3		1	1	0	0	
User 4	1	0	1	1	0	1
User 5	0		1	1		1
User 6	0	0	0	0	1	
User 7	1	0	1	1	0	1
User 8	0	1	1		0	1
User 9		0	1	1	1	1
User 10	1		0	0		0
User 11	0	1	1	1	0	1



Cómo entrenar un Auto Encoder

PASO 6

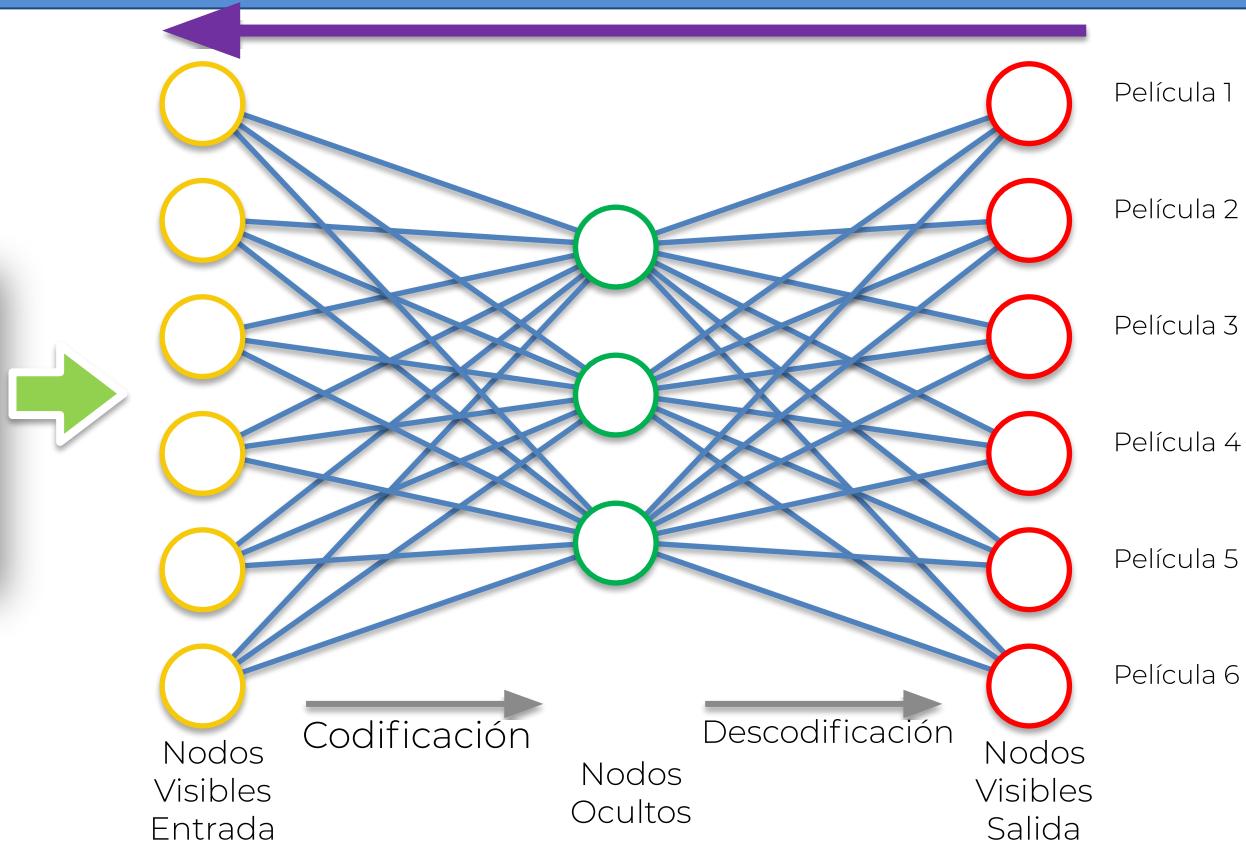
	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6
User 1	1	0				
User 2	0	1	0	0	1	0
User 3		1	1	0	0	
User 4	1	0	1	1	0	1
User 5	0		1	1		1
User 6	0	0	0	0	1	
User 7	1	0	1	1	0	1
User 8	0	1	1		0	1
User 9		0	1	1	1	1
User 10	1		0	0		0
User 11	0	1	1	1	0	1



Cómo entrenar un Auto Encoder

PASO 7

	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6
User 1	1	0				
User 2	0	1	0	0	1	0
User 3		1	1	0	0	
User 4	1	0	1	1	0	1
User 5	0		1	1		
User 6	0	0	0	0	1	
User 7	1	0	1	1	0	1
User 8	0	1	1		0	1
User 9		0	1	1	1	1
User 10	1		0	0		0
User 11	0	1	1	1	0	1

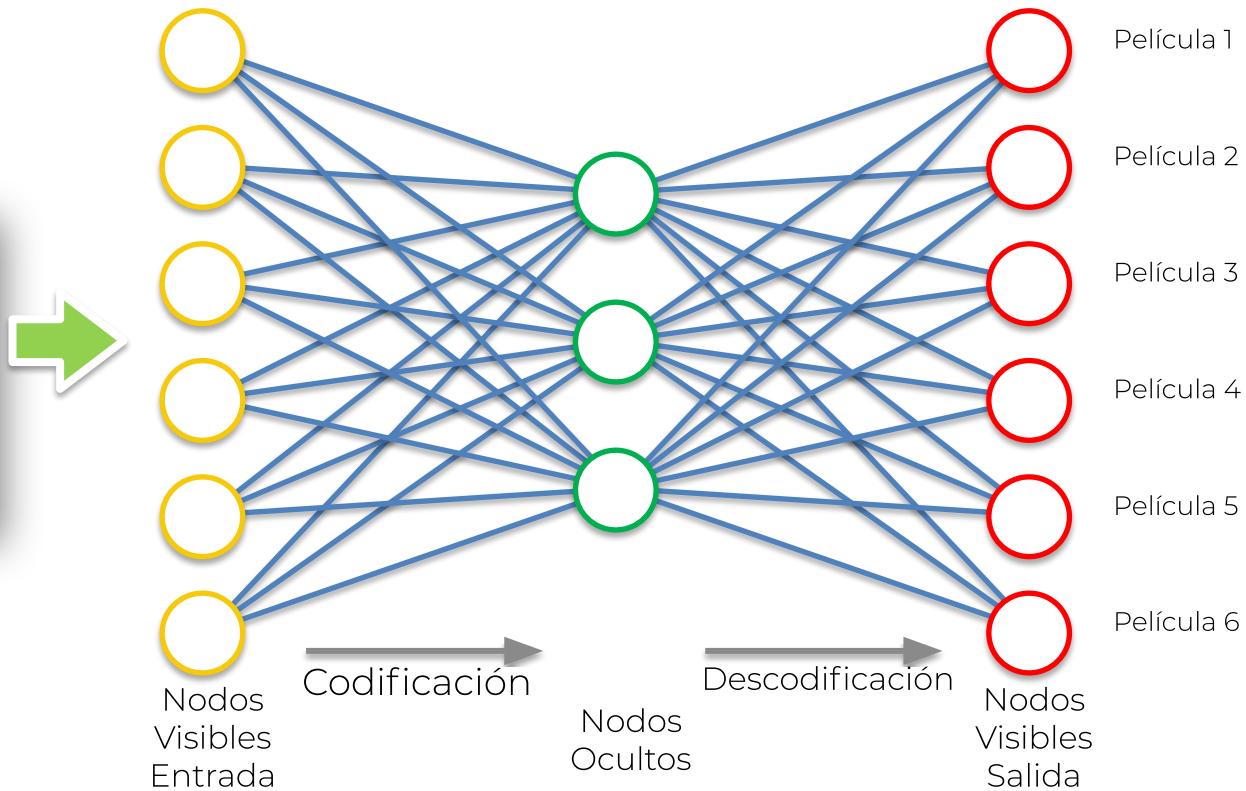


Cómo entrenar un Auto Encoder

PASO 8

	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6
User 1	1	0				
User 2	0	1	0	0	1	0
User 3		1	1	0	0	
User 4	1	0	1	1	0	1
User 5	0		1	1		1
User 6	0	0	0	0	1	
User 7	1	0	1	1	0	1
User 8	0	1	1		0	1
User 9		0	1	1	1	1
User 10	1		0	0		0
User 11	0	1	1	1	0	1

Repetimos



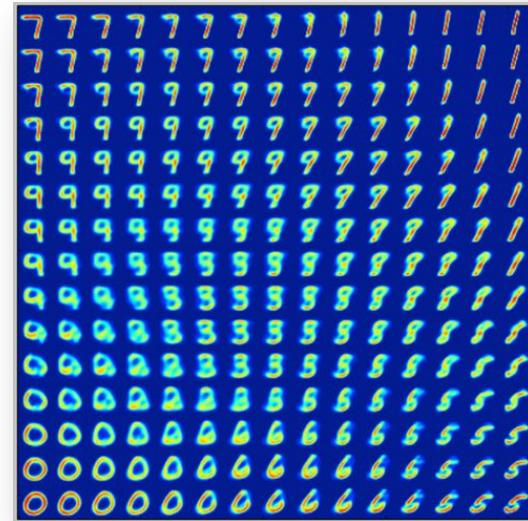
Cómo entrenar un Auto Encoder

Lectura adicional

Building Autoencoders in Keras
de Francois Chollet (2016)

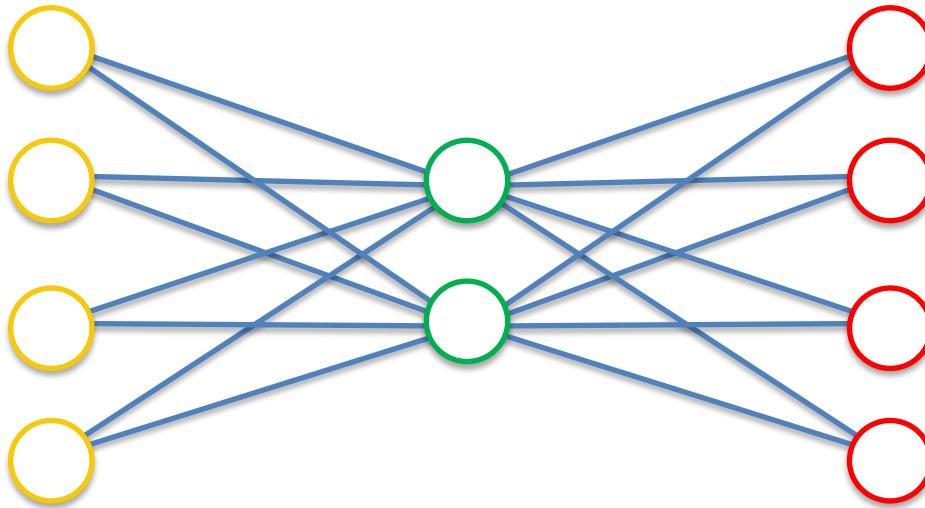
Link:

<https://blog.keras.io/building-autoencoders-in-keras.html>

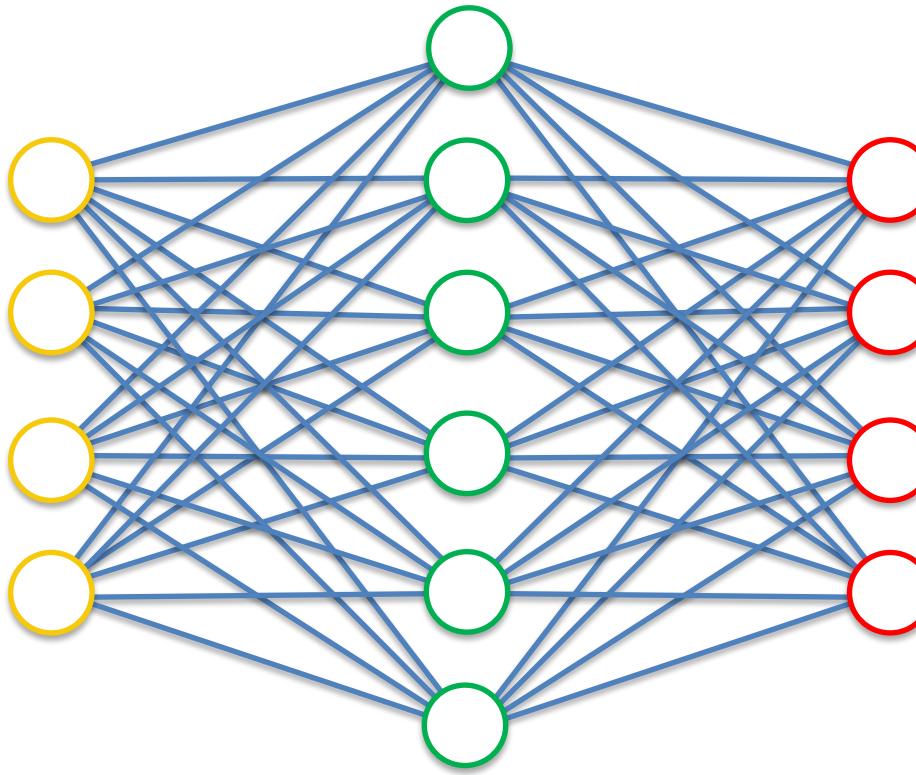


Sobrecompletar Capas Ocultas

Sobrecompletar Capas Ocultas

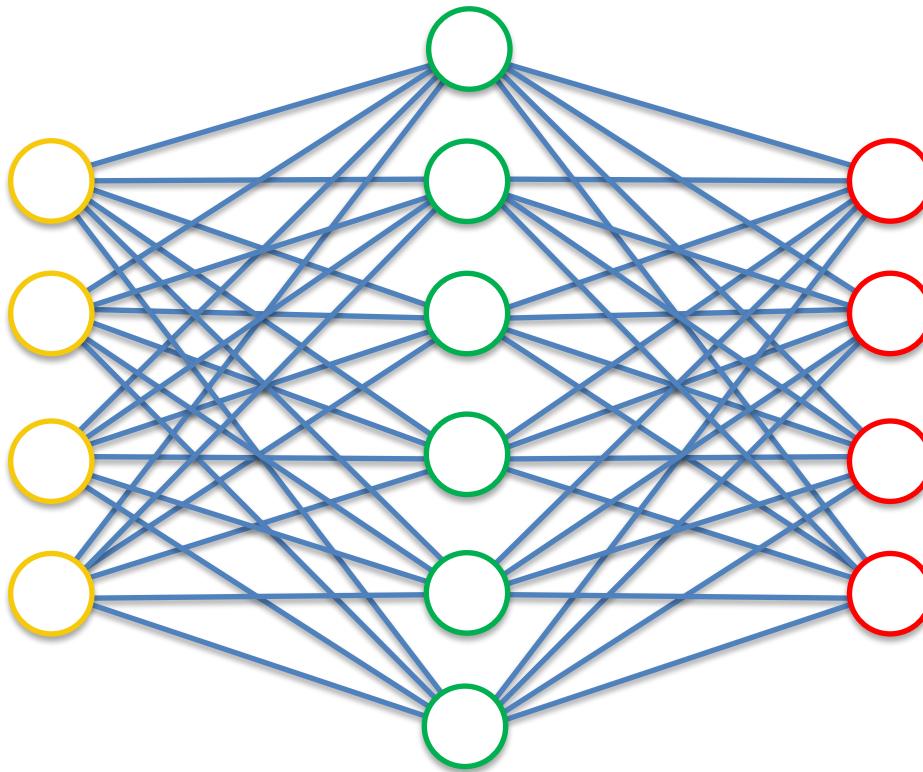


Sobrecompletar Capas Ocultas

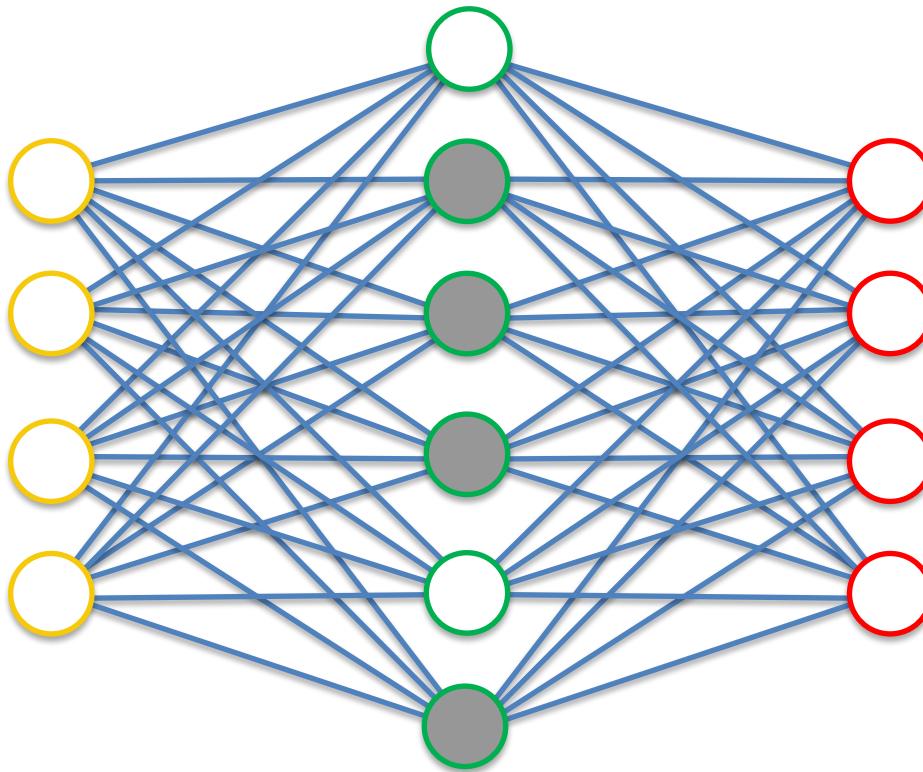


Sparse Autoencoders

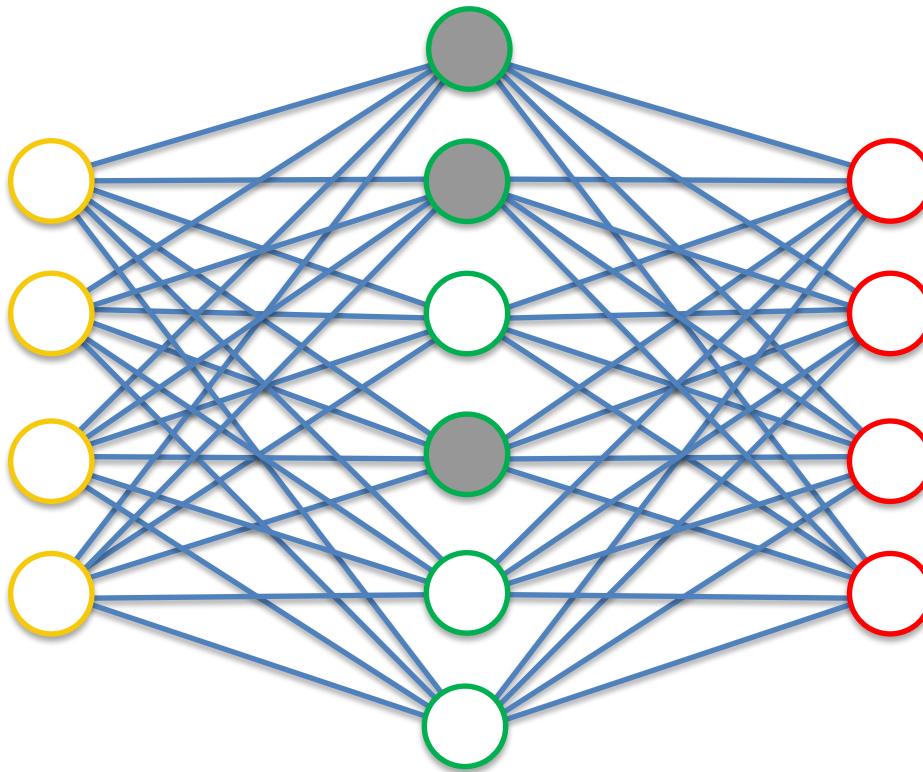
Sparse Autoencoders



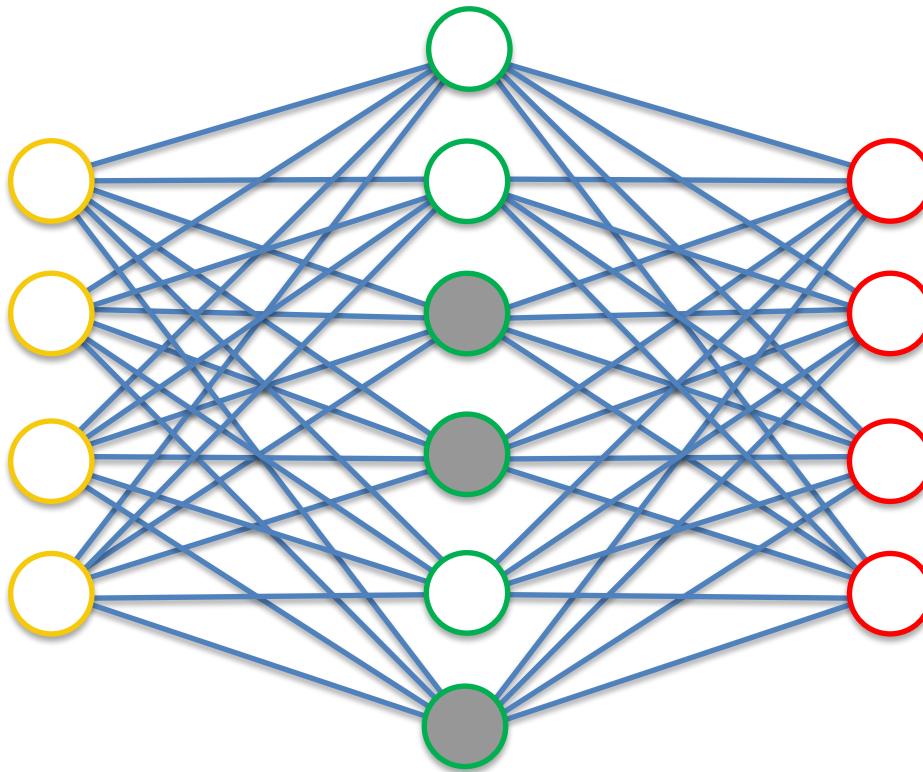
Sparse Autoencoders



Sparse Autoencoders



Sparse Autoencoders



Sparse Autoencoders

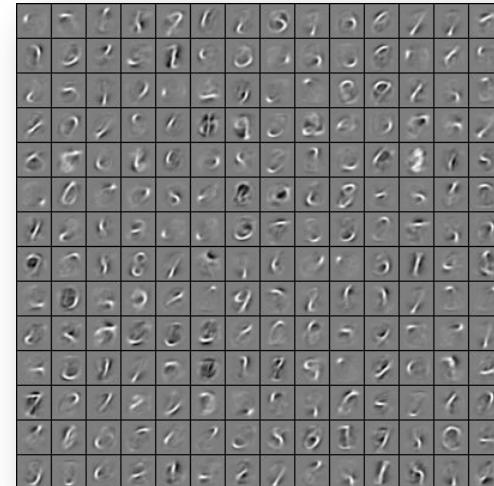
Lectura adicional

Deep Learning Tutorial - Sparse Autoencoder

de Chris McCormick (2014)

Link:

<http://mccormickml.com/2014/05/30/deep-learning-tutorial-sparse-autoencoder/>



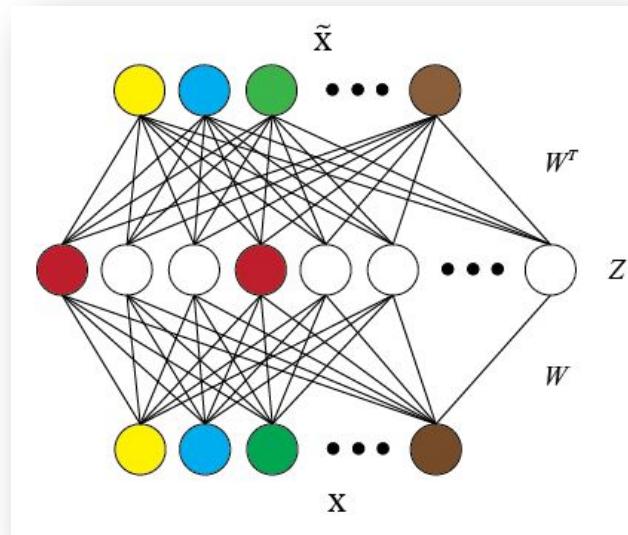
Sparse Autoencoders

Lectura adicional

Deep Learning: Sparse Autoencoders
de Eric Wilkinson (2014)

Link:

<http://www.erichwilkinson.com/blog/2014/11/19/deep-learning-sparse-autoencoders>



Sparse Autoencoders

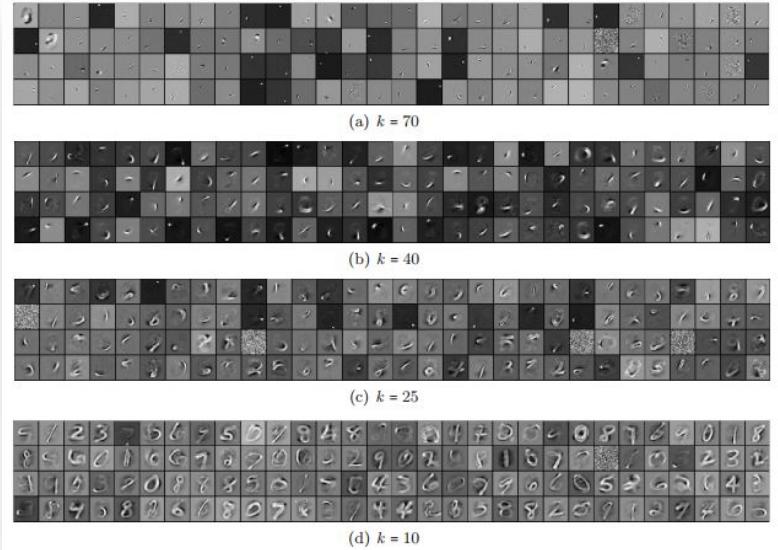
Lectura adicional

k-Sparse Autoencoders

de Alireza Makhzani et al. (2014)

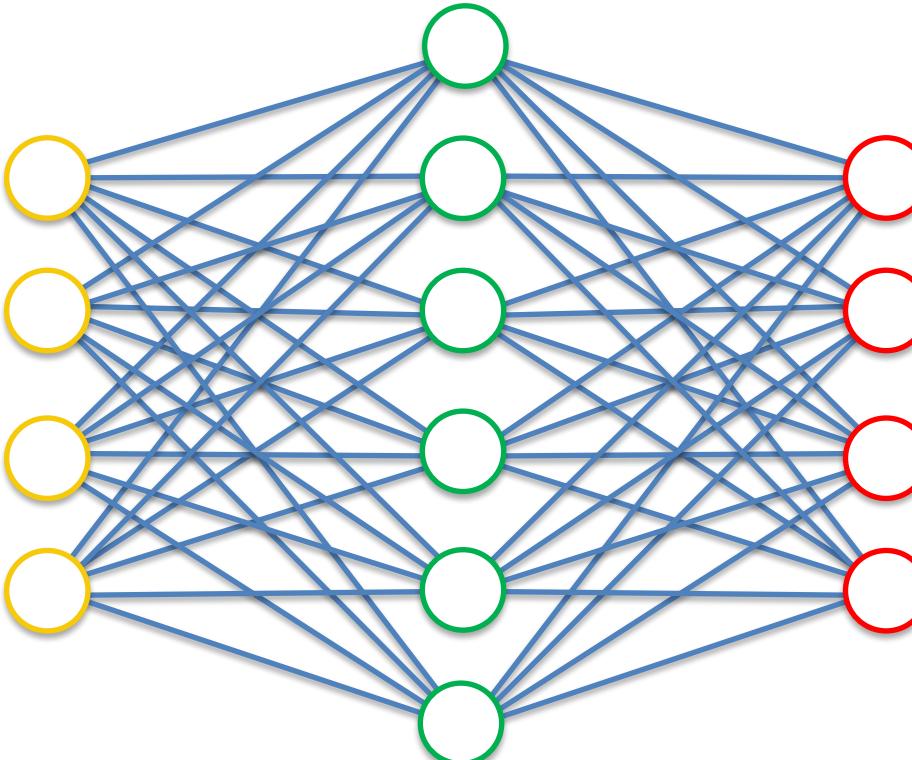
Link:

<https://arxiv.org/pdf/1312.5663.pdf>

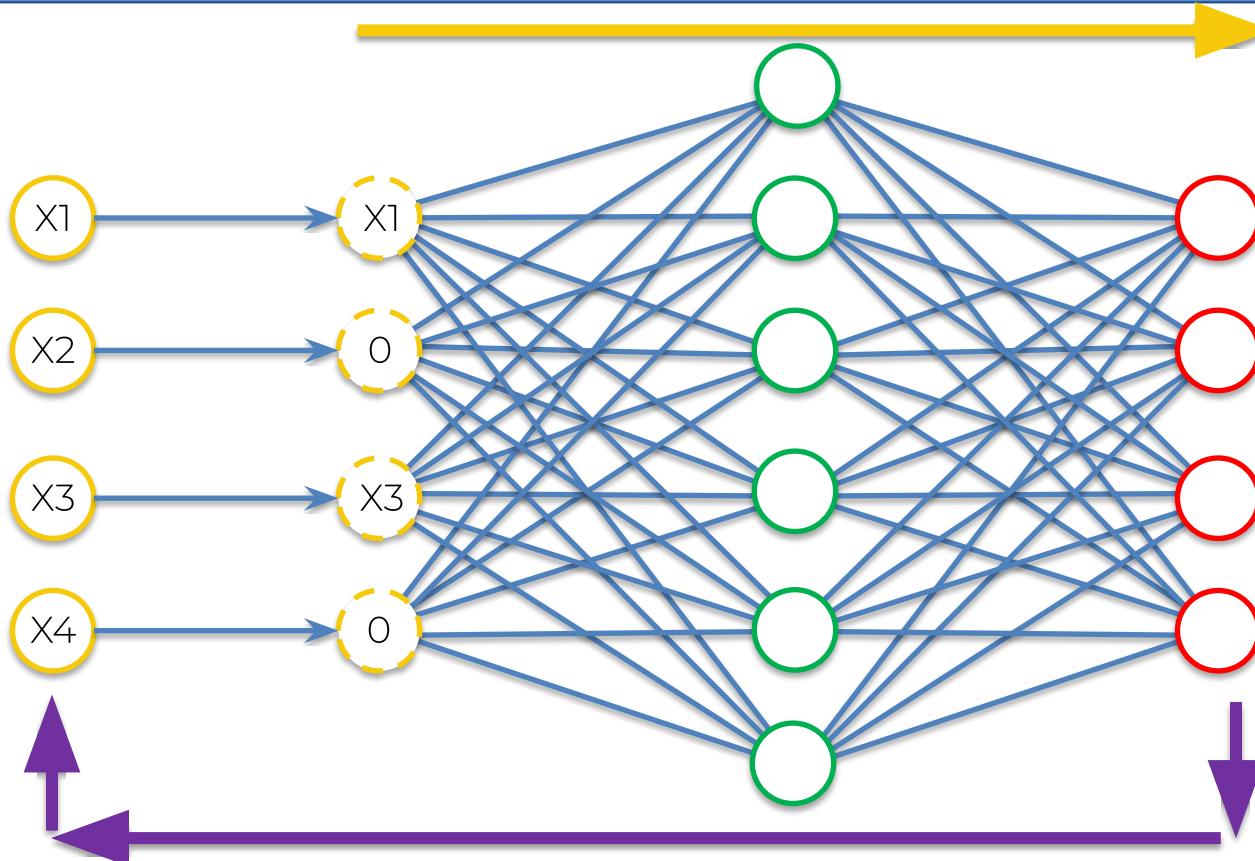


Denoising Autoencoders

Denoising Autoencoders



Denoising Autoencoders



Denoising Autoencoders

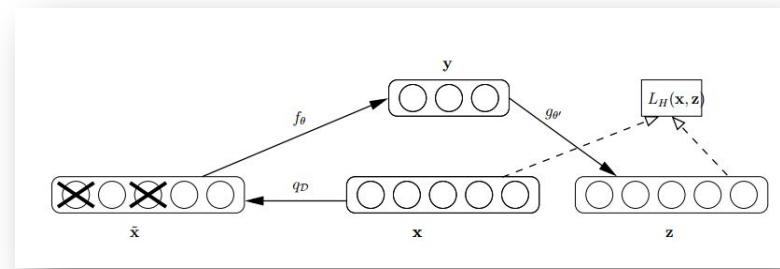
Lectura adicional

Extracting and Composing Robust Features with Denoising Autoencoders

de Pascal Vincent et al. (2008)

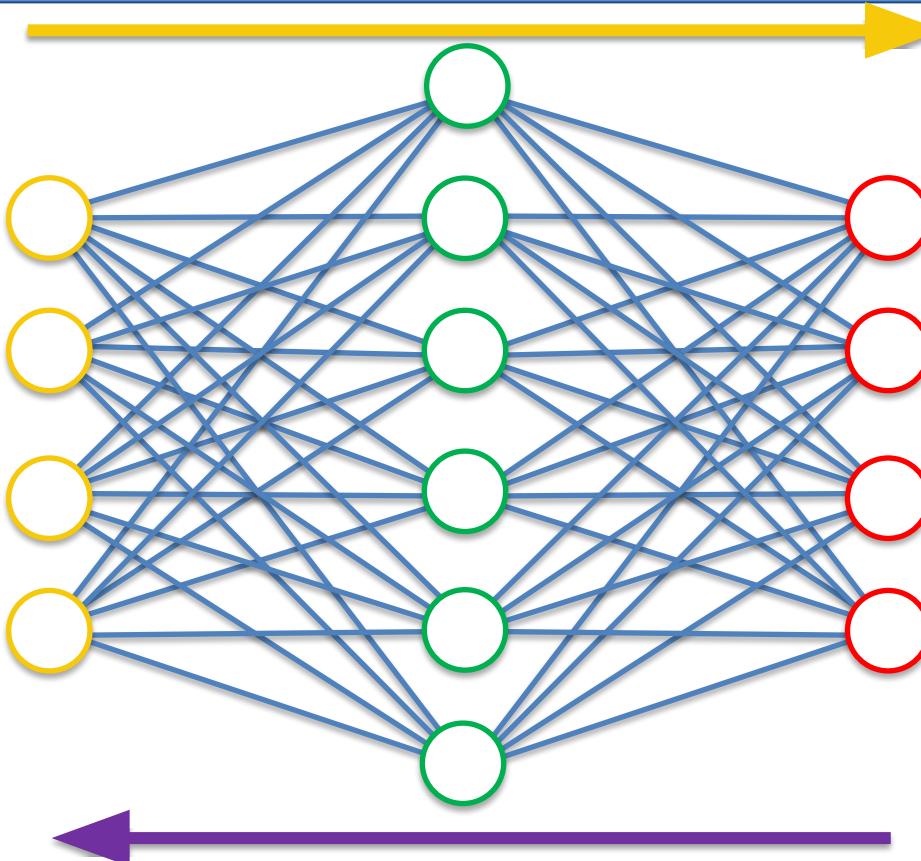
Link:

<http://www.cs.toronto.edu/~larocheh/publications/icml-2008-denoising-autoencoders.pdf>



Contractive Autoencoders

Contractive Autoencoders



Contractive Autoencoders

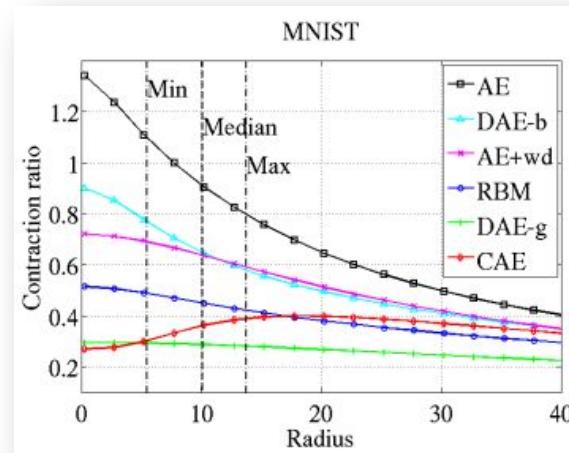
Lectura adicional

Contractive Auto-Encoders: Explicit Invariance During Feature Extraction

de Salah Rifai et al. (2011)

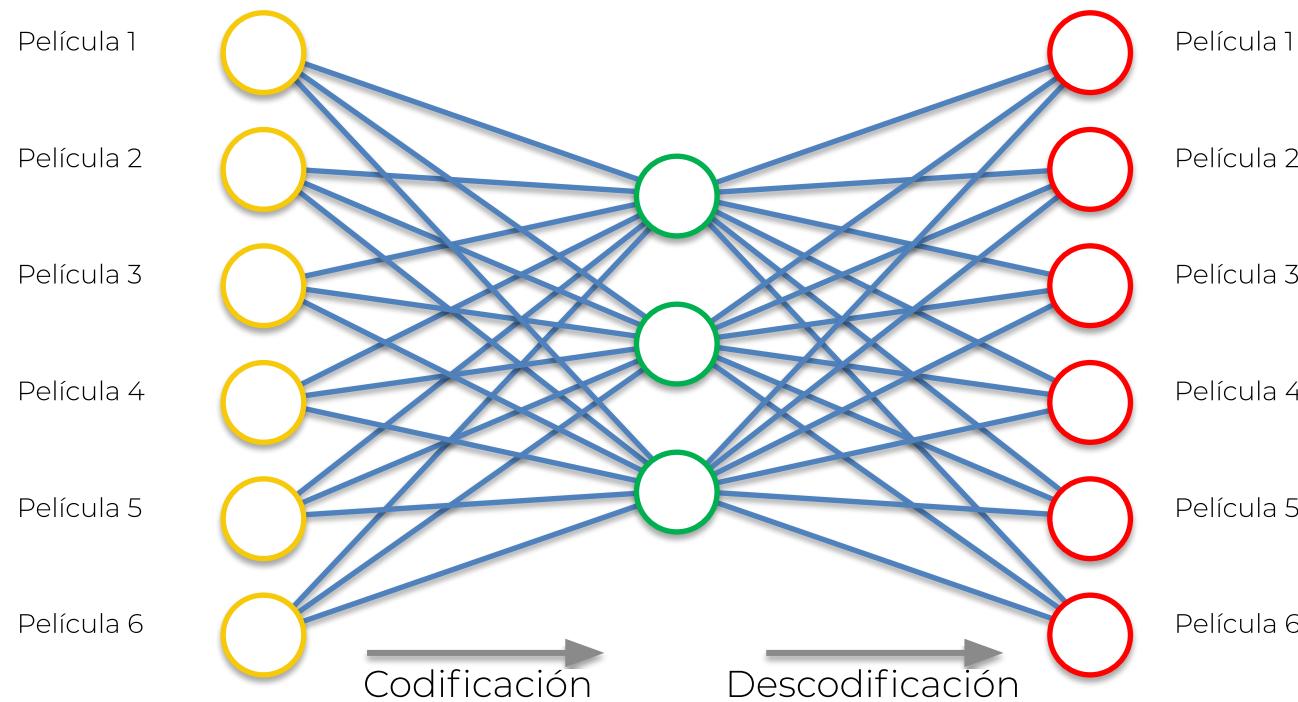
Link:

http://machinelearning.wustl.edu/mlpapers/paper_files/ICML2011Rifai_455.pdf

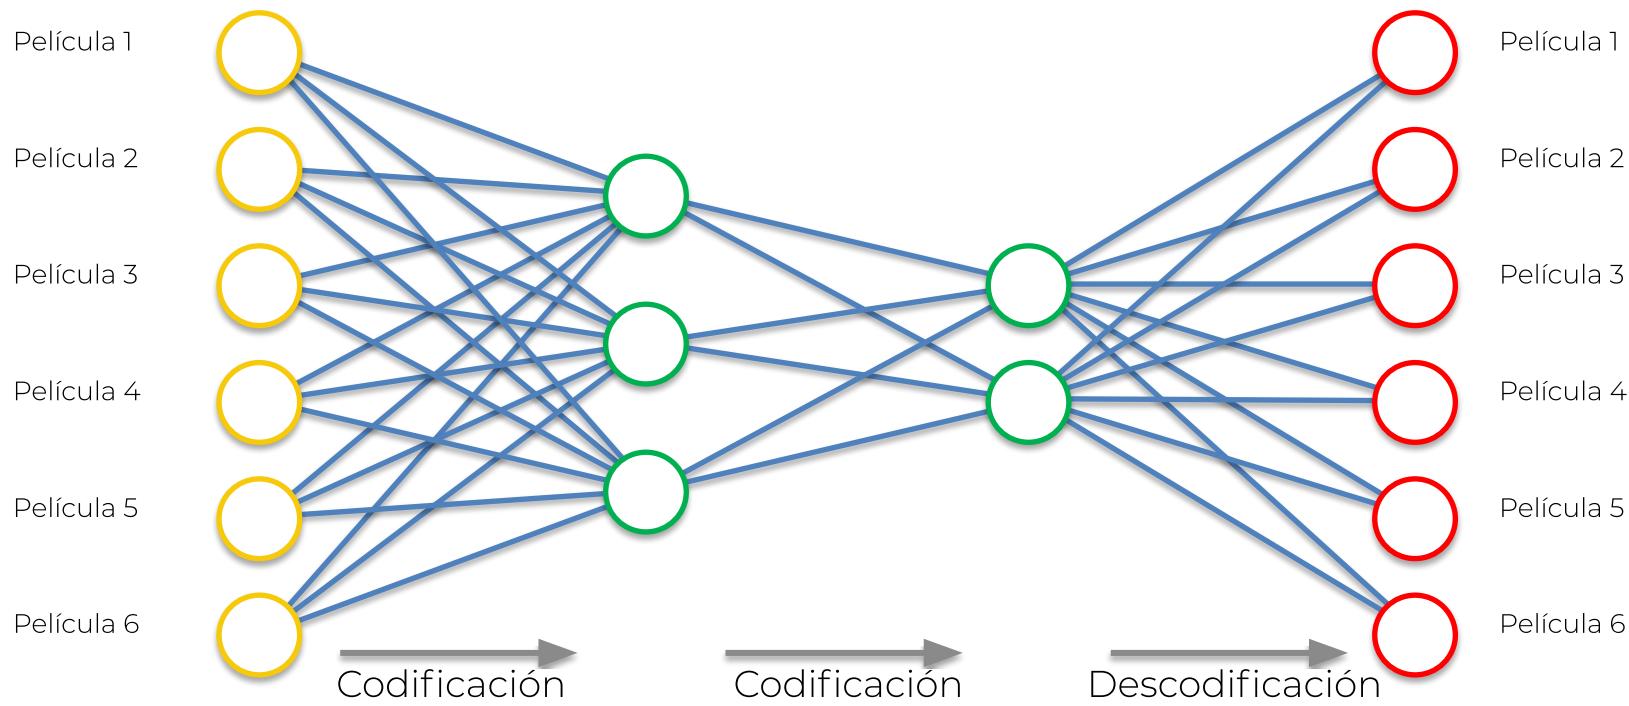


Stacked Autoencoders

Stacked Autoencoders



Stacked Autoencoders

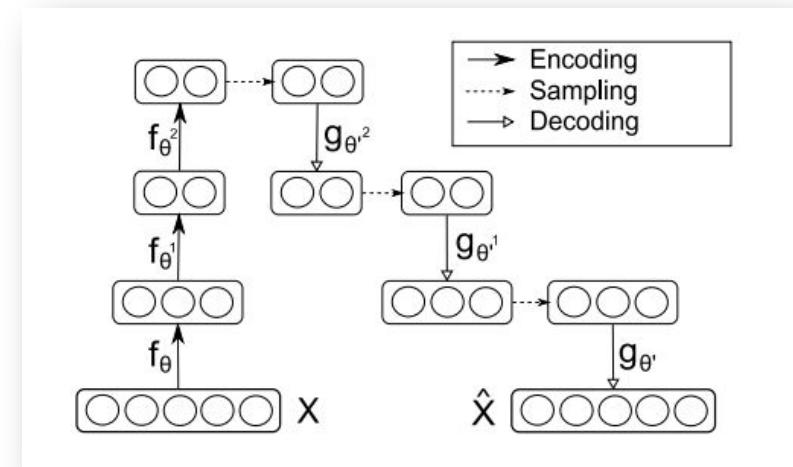


Stacked Autoencoders

Lectura adicional

*Stacked Denoising Autoencoders:
Learning Useful Representations in a
Deep Network with a Local Denoising
Criterion*

de Pascal Vincent et al. (2010)



Link:

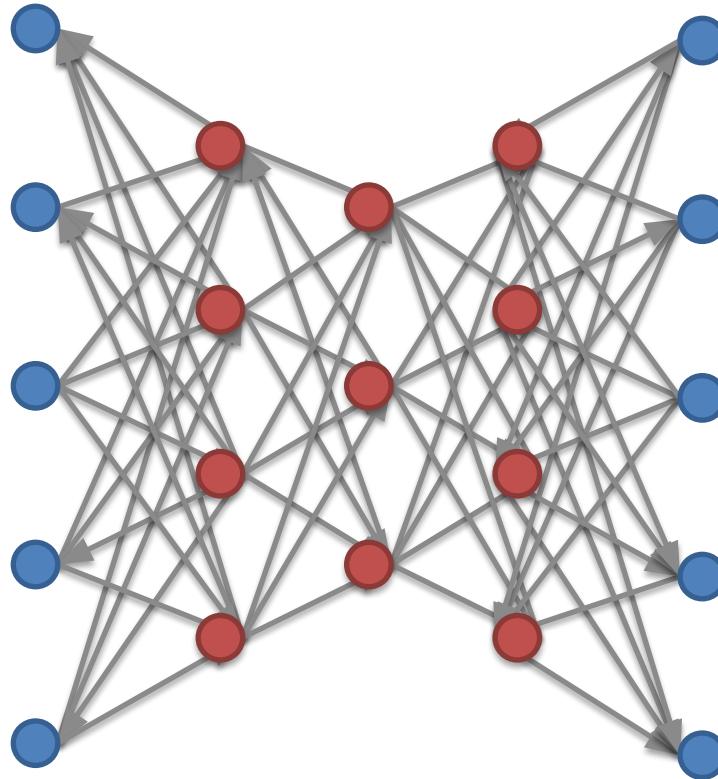
<http://www.jmlr.org/papers/volume11/vincent10a/vincent10a.pdf>

Deep Autoencoders

Deep Autoencoders

Stacked AE <> Deep AE

Deep Autoencoders



Deep Autoencoders

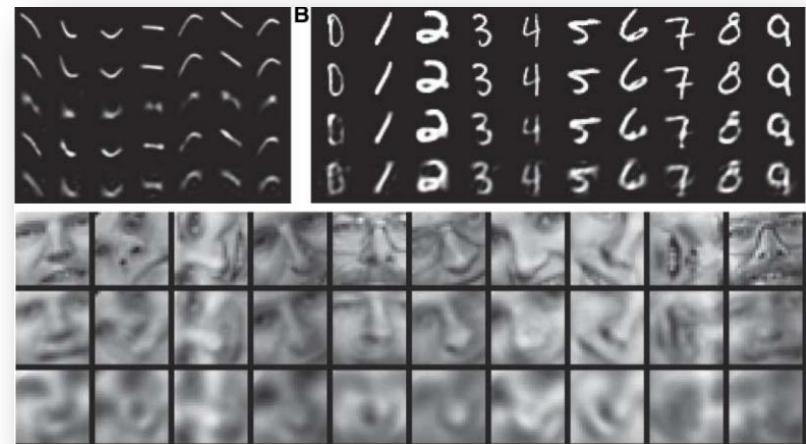
Lectura adicional

*Reducing the Dimensionality of Data
with Neural Networks*

de Geoffrey Hinton et al. (2006)

Link:

<https://www.cs.toronto.edu/~hinton/science.pdf>



EXTRA STUFF

***** EXTRA STUFF

<https://www.youtube.com/watch?v=b9COD6To5LM>

Unsupervised -> only the training data is available without any labels

“By restricting the number of hidden units we are forcing the network to learn a compressed representation of the input”

use | https://www.youtube.com/watch?v=b9COD6To5LM

Search

Autoencoders

- Tries to learn the identity function; unsupervised
- Not trivial if **constraints** are imposed
- Example (from Stanford.edu):
 - Given a 10x10 pixel image, we have a 100-vector
 - But if we have a 50-hidden unit layer, we are forced to learn a **compressed representation** of the input
 - I.i.d = difficult; Data in real world is has structure..

Full representation Compressed representation

Kin Gwn Lore

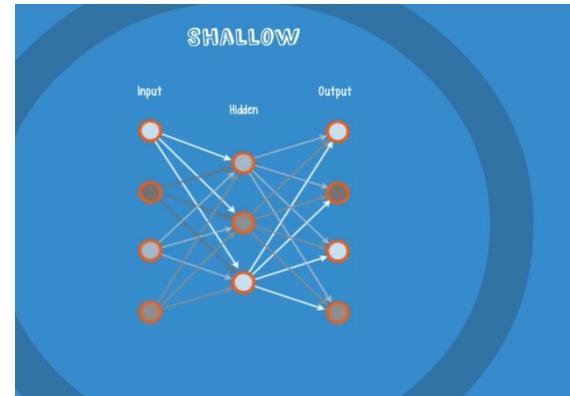
Self-aware learning laboratory

***** EXTRA STUFF

<https://www.youtube.com/watch?v=s96mYcicbpE>

~1m 10s – Wooow! “an rbm is an example of an autoencoder with only 2 layers”

Debatable...

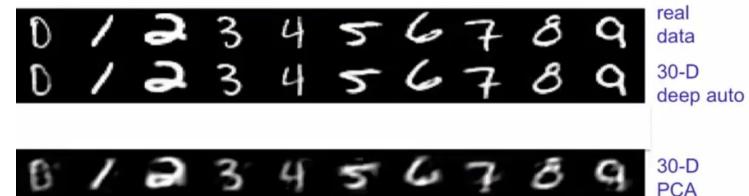


***** EXTRA STUFF

https://www.youtube.com/watch?v=_Ex1Ur85AVs

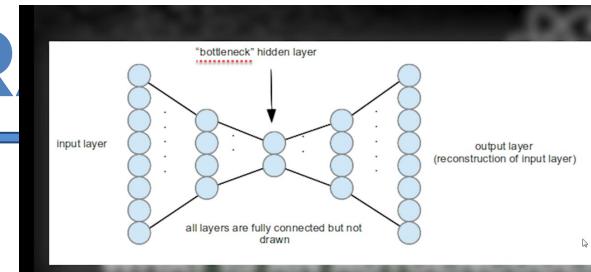
~3:30 DAE vs PCA

A comparison of methods for compressing digit images to 30 real numbers



EXTR

<https://www.youtube.com/watch?v=Rdpbnd0pCil>



***** EXTRA STUFF

<https://blog.keras.io/building-autoencoders-in-keras.html>

The Keras Blog

Keras is a Deep Learning library for Python, that is simple, modular, and extensible.

Archives GitHub Documentation Google Group

Building Autoencoders in Keras

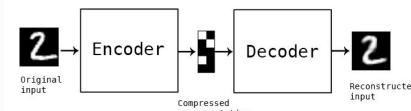
In this tutorial, we will answer some common questions about autoencoders, and we will cover code examples of the following models:

- a simple autoencoder based on a fully-connected layer
- a sparse autoencoder
- a deep fully-connected autoencoder
- a deep convolutional autoencoder
- an image denoising model
- a sequence-to-sequence autoencoder
- a variational autoencoder

Sat 14 May 2016
By [François Chollet](#)
In [Tutorials](#).

Note: all code examples have been updated to the Keras 2.0 API on March 14, 2017. You will need Keras version 2.0.0 or higher to run them.

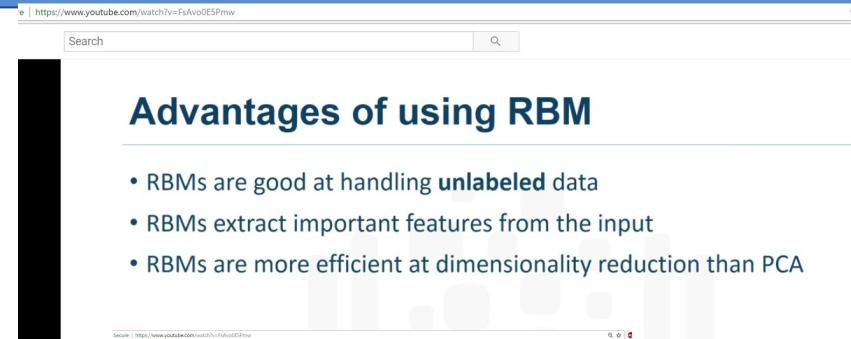
What are autoencoders?



"Autoencoding" is a data compression algorithm where the compression and decompression functions are 1) data-specific, 2) lossy, and 3) learned automatically from examples rather than engineered by a human. Additionally, in almost all contexts where the term "autoencoder" is used, the compression and decompression functions are implemented with neural networks.

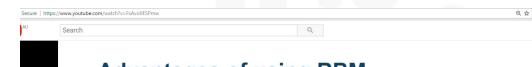
***** EXTRA STUFF

<https://www.youtube.com/watch?v=FsAvo0E5Pmw>



Advantages of using RBM

- RBMs are good at handling **unlabeled** data
- RBMs extract important features from the input
- RBMs are more efficient at dimensionality reduction than PCA



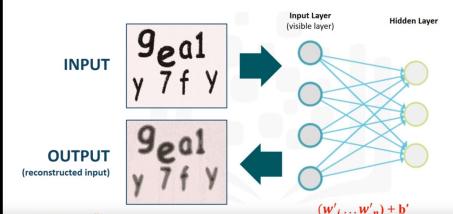
Advantages of using RBM

Note that the Restricted Boltzmann Machine is a type of

AUTOENCODER



Learning Process of RBMs



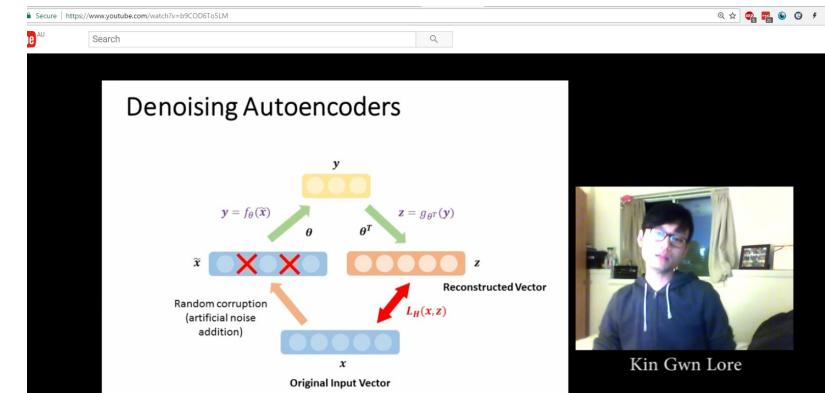
***** EXTRA STUFF

<https://www.youtube.com/watch?v=b9COD6To5LM>

Denoising autoencoders ~2:30

“Randomly corrupt the input data”

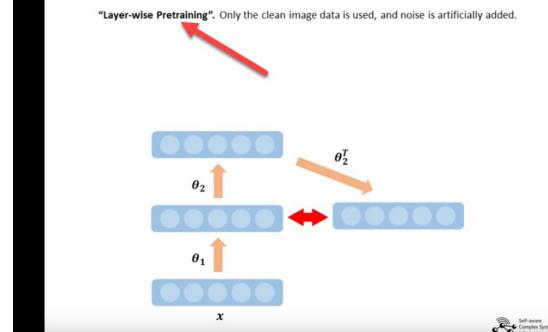
Layer-wise pre-training ~3:40



Kin Gwn Lore

Stacked Denoising Autoencoders (SDA)

“Layer-wise Pretraining”. Only the clean image data is used, and noise is artificially added.

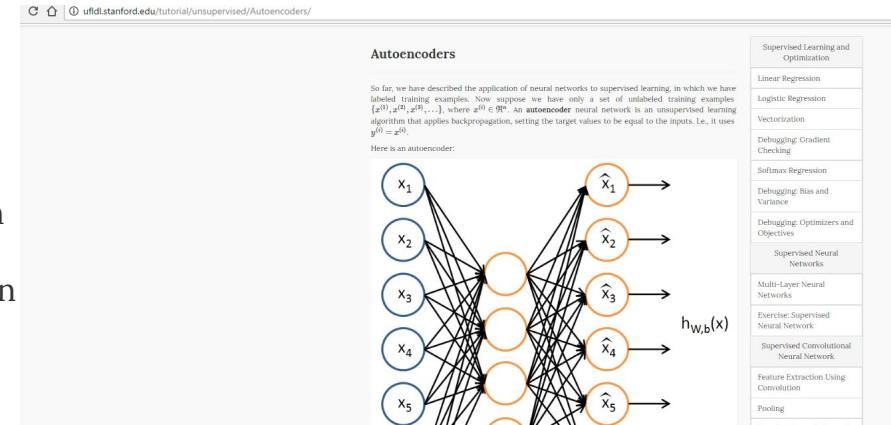


Kin Gwn Lore

***** EXTRA STUFF

<http://ufldl.stanford.edu/tutorial/unsupervised/Autoencoders/>

If the input were completely random—say, each x_i comes from an IID Gaussian independent of the other features—then this compression task would be very difficult. But if there is structure in the data, for example, if some of the input features are correlated, then this algorithm will be able to discover some of those correlations.



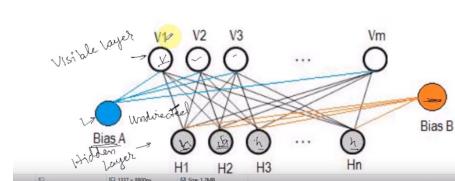
BTW: If you are using a tanh activation function, then we think of a neuron as being inactive when it outputs values close to -1.

Sparsity: Our argument above relied on the number of hidden units s_2 being small. But even when the number of hidden units is large (perhaps even greater than the number of input pixels), we can still discover interesting structure, by imposing other constraints on the network. In particular, if we impose a "sparsity" constraint on the hidden units, then the autoencoder will still discover interesting structure in the data, even if the number of hidden units is large. <<! Find out more info about this.

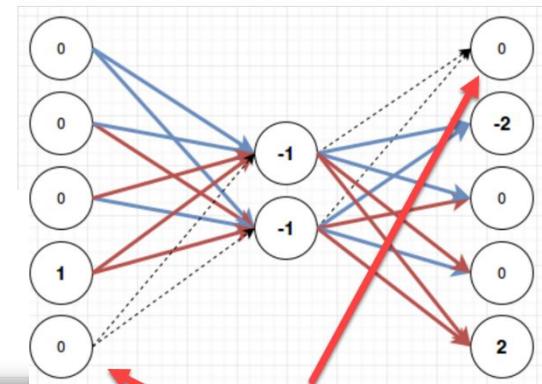
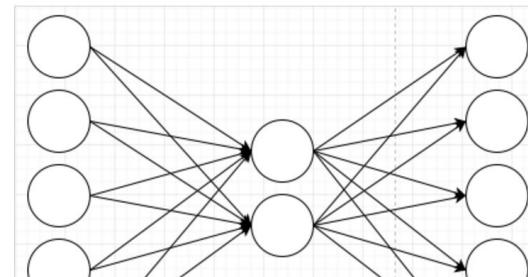
***** EXTRA STUFF

<https://probablydance.com/2016/04/30/neural-networks-are-impressively-good-at-compression/>

(not obvious, start with this): If you count the number of connections on that last picture you will notice that there are fewer than there were in the first network. There were $5 \times 5 = 25$ at first, now there are $5 \times 2 \times 2 = 20$. That reduction would be larger if I had more input and output nodes. For 10 nodes that reduction would be from 100 connections down to 40 when inserting two nodes in the middle.



Secure <https://probablydance.com/2016/04/30/neural-networks-are-impressively-good-at-compression/>
A simple network like the one above can learn simple patterns like this. If we want t
we have to create a network that has more layers. Let's insert a layer with two neuror



it is the
res exp

***** EXTRA STUFF

<https://probablydance.com/2016/04/30/neural-networks-are-impressively-good-at-compression/>
Interesting:

1. Explanation why for Tanh(): -1 = non-active and 1= active: because after that a softmax function is used and it's e to the power of something. "Also by using an exponent I get no negative numbers. 2^{-96} is just a number that's very close to zero." -> we need to use softmax
2. Explanation why we need Tanh(): "In fact since I just used linear math in my explanation above and didn't need to run tanh() on anything, couldn't softmax have learned those weights even without using a tanh layer? In theory it could, but in practice it will keep on bumping up the numbers to get small improvements and you will very quickly overflow floating point numbers. The tanh layer in the middle is then just responsible for keeping the numbers small: it clamps its outputs to the range -1 to 1, and the inputs won't grow past a certain point either because at some point a bigger number just means that the output goes from 0.995 to 0.997. But since they can usually still get a small improvement you don't lose that nice property of softmax where neurons can keep on edging out small improvements over each other." -> we need Tanh()

Summary: "So if you just use softmax your weights will explode and if you just use tanh your weights will stagnate too soon before a good solution emerges. If you use both you get nice greedy behavior where the connections keep on looking for better solutions, but you also keep the weights from exploding." -> we need both!

***** EXTRA STUFF

<https://probablydance.com/2016/04/30/neural-networks-are-impressively-good-at-compression/>

Interesting:

Now all I've shown is that my network can learn the rule that when input n fires, it needs to fire output n+1. It should be easy to see that just by creating a different permutation of the weights of the connections I can teach my network that for any input neuron x, it should fire any other output neuron y. The impressive part is that without the middle layer I would have needed $11 \times 11 = 121$ connections to have a network that can learn any combination, and with that middle layer I can do it in only $11 \times 2 \times 2 = 44$ connections.

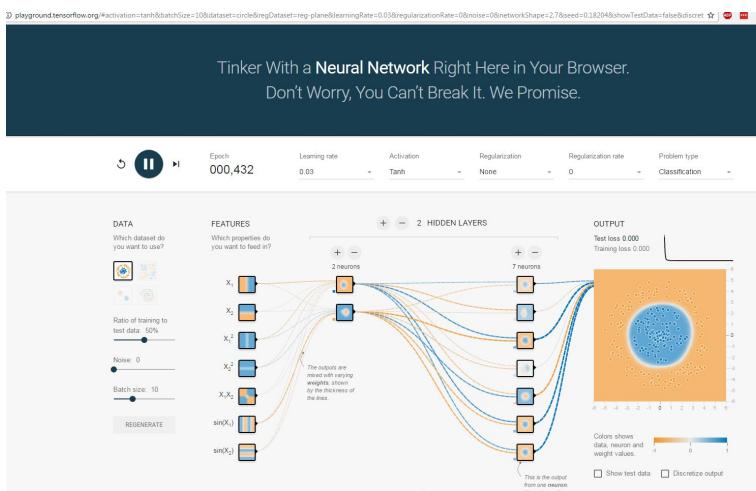
Eleven inputs and outputs is close to the limit for what you can do with two nodes in the middle. If you ask it to do much more the weights will fluctuate and they will keep on PASOping on each others toes. But with three nodes in the middle I can learn these patterns for something like thirty inputs and outputs, and with four nodes, I can learn direct connections for more than a hundred inputs and outputs. It doesn't grow linearly because the number of combinations doesn't grow linearly.

But backprop DOES scale with the #of connections. So we win! "Luckily the back propagation algorithm scales with the number of connections, not with the number of combinations. So with a linear increase in computation cost I get a super-linear increase in the amount of stuff I can learn."

***** EXTRA STUFF

Understand how to use this

<http://playground.tensorflow.org>



***** EXTRA STUFF

<https://www.microsoft.com/en-us/research/publication/three-classes-of-deep-learning-architectures-and-their-applications-a-tutorial-survey/?from=http%3A%2F%2Fresearch.microsoft.com%2Fpubs%2F192937%2Ftransactions-apsipa.pdf>

<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/Transactions-APSI PA.pdf>

The screenshot shows a Microsoft Word document with the title 'APSI PA.pdf'. At the top right, it says 'Secure | https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/Transactions-APSI PA.pdf' and '1 / 28'. The main content is a paper titled 'Three Classes of Deep Learning Architectures and Their Applications: A Tutorial Survey' by Li Deng. The abstract discusses the history of deep learning, its connection to representation learning, and provides a taxonomy of deep learning architectures. A red arrow points to the author's name, Li Deng, and his affiliation: 'Microsoft Research, Redmond, WA 98052, USA'.

Three Classes of Deep Learning Architectures and Their Applications: A Tutorial Survey

Li Deng
Microsoft Research, Redmond, WA 98052, USA
Li.Deng@microsoft.com, Tel: 425-706-2719

Abstract— In this invited paper, my overview material on the same topic as presented in the plenary overview session of APSIPA-2011 and the tutorial material presented in the same conference (Deng, 2011) are expanded and updated to include more recent developments in deep learning. The previous and the updated materials cover both theory and applications, and analyze its future directions. The goal of this tutorial survey is to introduce the emerging area of deep learning or hierarchical learning to the APSIPA community. Deep learning refers to a class of machine learning techniques, developed largely since 2006, where many stages of nonlinear information processing in hierarchical architectures are exploited for pattern classification and for feature learning. In the more recent literature, it is also connected to representation learning, which involves a hierarchy of features or concepts where higher-level concepts are defined from lower-level ones and where the same lower-level concepts help to define higher-level ones. In this tutorial, a brief history of deep learning research is discussed first. Then, a classificatory scheme is developed to analyze and summarize major work reported in the deep learning literature. Using this scheme, I provide a taxonomy-oriented survey on the existing deep architectures and algorithms in the literature, and categorize

***** **EXTRA STUFF**

AE vs RBM

<https://www.quora.com/What-is-the-difference-between-Deep-Boltzmann-Machine-Deep-Belief-Networks-and-Deep-Auto-Encoders-Is-there-is-tutorial-which-explains-the-difference-between-procedures-for-pre-training-and-fine-tuning-the-above-networks-pseudo-code-would-be-of-additional-help>